

## Hands-on Workshop: Deep Learning for Computer Vision

March 2019

Instructor: Santosh Chapaneri

- The Fashion MNIST database has a database of fashion accessories.
- Fashion-MNIST: a Novel Image Dataset for Benchmarking Machine Learning Algorithms. Han Xiao, Kashif Rasul, Roland Vollgraf. arXiv:1708.07747, 2017.
- The training set has 60,000 samples. The test set has 10,000 samples.
- The fashion accessories are size-normalized and centered in a fixed-size image.
- We will train Multi-layer Perceptron, Deep Multi-layer Perceptron and CNN classifier using Keras for Fashion MNIST dataset.

```
In [1]: import numpy as np
import pandas as pd
import tensorflow as tf
import keras
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from keras.utils import to_categorical
```

Using TensorFlow backend.

```
In [2]: # Load the fashion-mnist pre-shuffled train data and test data
(X_train, Y_train), (X_test, Y_test) = tf.keras.datasets.fashion_mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
32768/29515 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
26427392/26421880 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
8192/5148 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
4423680/4422102 [=====] - 0s 0us/step
```

```
In [3]: # Print training set shape - note there are 60,000 training data of image size of 28x28, 60,000 train labels)
print("X_train shape:", X_train.shape, "Y_train shape:", Y_train.shape)

X_train shape: (60000, 28, 28) Y_train shape: (60000,)
```

## Visualize the data

```
In [4]: # Define the Labels
fashion_mnist_labels = ["T-shirt/top", # index 0
                        "Trouser",      # index 1
                        "Pullover",     # index 2
                        "Dress",         # index 3
                        "Coat",          # index 4
                        "Sandal",        # index 5
                        "Shirt",         # index 6
                        "Sneaker",       # index 7
                        "Bag",           # index 8
                        "Ankle boot"]    # index 9

# Image index, you can pick any number between 0 and 59,999
img_index = 5

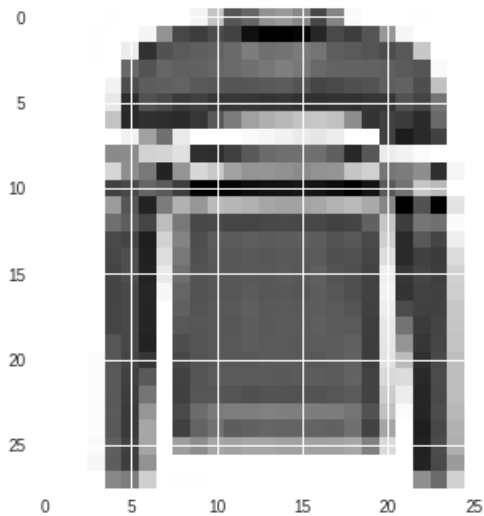
# y_train contains the labels, ranging from 0 to 9
label_index = Y_train[img_index]

# Print the label, for example 2 Pullover
print ("Y = " + str(label_index) + " " + (fashion_mnist_labels[label_index]))

# # Show one of the images from the training dataset
plt.imshow(X_train[img_index])

Y = 2 Pullover
```

Out[4]: <matplotlib.image.AxesImage at 0x7f0d6e51c748>



```
In [0]: img_rows, img_cols = 28, 28

# MLP
X_train_mlp = X_train.reshape(X_train.shape[0],img_rows*img_cols)
Y_train_mlp = Y_train
X_test_mlp = X_test.reshape(X_test.shape[0],img_rows*img_cols)
Y_test_mlp = Y_test

# CNN
X_train_cnn = X_train.reshape(X_train.shape[0],img_rows,img_cols,1)
Y_train_cnn = Y_train
X_test_cnn = X_test.reshape(X_test.shape[0],img_rows,img_cols,1)
Y_test_cnn = Y_test
```

```
In [6]: print(X_train_mlp.shape)
print(X_train_cnn.shape)

(60000, 784)
(60000, 28, 28, 1)
```

```
In [0]: X_train_mlp = X_train_mlp.astype('float32')
X_test_mlp = X_test_mlp.astype('float32')
X_train_mlp /= 255
X_test_mlp /= 255

X_train_cnn = X_train_cnn.astype('float32')
X_test_cnn = X_test_cnn.astype('float32')
X_train_cnn /= 255
X_test_cnn /= 255
```

```
In [0]: # Convert class vectors to binary class matrices
num_classes = 10

Y_train_mlp = keras.utils.to_categorical(Y_train_mlp, num_classes)
Y_test_mlp = keras.utils.to_categorical(Y_test_mlp, num_classes)

Y_train_cnn = keras.utils.to_categorical(Y_train_cnn, num_classes)
Y_test_cnn = keras.utils.to_categorical(Y_test_cnn, num_classes)
```

```
In [9]: Y_train_cnn[:5,:]
```

```
Out[9]: array([[0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],
               [1., 0., 0., 0., 0., 0., 0., 0., 0., 0.],
               [1., 0., 0., 0., 0., 0., 0., 0., 0., 0.],
               [0., 0., 0., 1., 0., 0., 0., 0., 0., 0.],
               [1., 0., 0., 0., 0., 0., 0., 0., 0., 0.]], dtype=float32)
```

```
In [0]: # Split data to optimize classifier during training
X_train_mlp, X_val_mlp, Y_train_mlp, Y_val_mlp = train_test_split(X_train_mlp,
                                                                    Y_train_mlp,
                                                                    test_size=0.2)

X_train_cnn, X_val_cnn, Y_train_cnn, Y_val_cnn = train_test_split(X_train_cnn,
                                                                    Y_train_cnn,
                                                                    test_size=0.2)
```

```
In [11]: print(X_train_mlp.shape)
print(X_val_mlp.shape)

print(X_train_cnn.shape)
print(X_val_cnn.shape)

(48000, 784)
(12000, 784)
(48000, 28, 28, 1)
(12000, 28, 28, 1)
```

## Multi Layer Perceptron

```
In [0]: from keras.models import Sequential
from keras.layers import Dense
from keras.optimizers import SGD
from keras.datasets import mnist
from keras.utils import np_utils
```

In [13]: *# Multilayer Perceptron model*

```
batch_size = 256
num_epochs = 50

model = Sequential()

model.add(Dense(input_dim=784, activation='sigmoid',
                 units=625, kernel_initializer='normal'))

model.add(Dense(input_dim=625, activation='softmax',
                 units=10, kernel_initializer='normal'))

model.compile(optimizer=SGD(lr=0.05),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/framework/op\_def\_library.py:263: colocate\_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 625)	490625
dense_2 (Dense)	(None, 10)	6260
Total params: 496,885		
Trainable params: 496,885		
Non-trainable params: 0		

```
In [14]: history = model.fit(X_train_mlp, Y_train_mlp,  
                             batch_size = batch_size,  
                             epochs = num_epochs,  
                             verbose = 1,  
                             validation_data = (X_val_mlp, Y_val_mlp))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math\_ops.py:3066: to\_int32 (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 48000 samples, validate on 12000 samples

Epoch 1/50

48000/48000 [=====] - 4s 75us/step - loss: 1.3990 - acc: 0.6144 - val\_loss: 0.9937 - val\_acc: 0.7046

Epoch 2/50

48000/48000 [=====] - 1s 21us/step - loss: 0.8636 - acc: 0.7324 - val\_loss: 0.7843 - val\_acc: 0.7523

Epoch 3/50

48000/48000 [=====] - 1s 20us/step - loss: 0.7346 - acc: 0.7546 - val\_loss: 0.7099 - val\_acc: 0.7500

Epoch 4/50

48000/48000 [=====] - 1s 20us/step - loss: 0.6726 - acc: 0.7682 - val\_loss: 0.6556 - val\_acc: 0.7751

Epoch 5/50

48000/48000 [=====] - 1s 22us/step - loss: 0.6325 - acc: 0.7810 - val\_loss: 0.6275 - val\_acc: 0.7865

Epoch 6/50

48000/48000 [=====] - 1s 21us/step - loss: 0.6035 - acc: 0.7891 - val\_loss: 0.6007 - val\_acc: 0.7917

Epoch 7/50

48000/48000 [=====] - 1s 20us/step - loss: 0.5812 - acc: 0.7975 - val\_loss: 0.5873 - val\_acc: 0.7985

Epoch 8/50

48000/48000 [=====] - 1s 20us/step - loss: 0.5636 - acc: 0.8042 - val\_loss: 0.5678 - val\_acc: 0.8022

Epoch 9/50

48000/48000 [=====] - 1s 21us/step - loss: 0.5478 - acc: 0.8093 - val\_loss: 0.5555 - val\_acc: 0.8078

Epoch 10/50

48000/48000 [=====] - 1s 21us/step - loss: 0.5351 - acc: 0.8136 - val\_loss: 0.5463 - val\_acc: 0.8105

Epoch 11/50

48000/48000 [=====] - 1s 20us/step - loss: 0.5242 - acc: 0.8172 - val\_loss: 0.5323 - val\_acc: 0.8153

Epoch 12/50

48000/48000 [=====] - 1s 21us/step - loss: 0.5157 - acc: 0.8200 - val\_loss: 0.5211 - val\_acc: 0.8212

Epoch 13/50

48000/48000 [=====] - 1s 21us/step - loss: 0.5077 - acc: 0.8228 - val\_loss: 0.5156 - val\_acc: 0.8193

Epoch 14/50

48000/48000 [=====] - 1s 21us/step - loss: 0.5002 - acc: 0.8249 - val\_loss: 0.5198 - val\_acc: 0.8162

Epoch 15/50

48000/48000 [=====] - 1s 21us/step - loss: 0.4940 - acc: 0.8273 - val\_loss: 0.5021 - val\_acc: 0.8285

Epoch 16/50

48000/48000 [=====] - 1s 21us/step - loss: 0.4880 - acc: 0.8286 - val\_loss: 0.5011 - val\_acc: 0.8268

Epoch 17/50

48000/48000 [=====] - 1s 21us/step - loss: 0.4828 - acc: 0.8315 - val\_loss: 0.5003 - val\_acc: 0.8242

Epoch 18/50

48000/48000 [=====] - 1s 21us/step - loss: 0.4779 - acc: 0.8328 - val\_loss: 0.4903 - val\_acc: 0.8293

Epoch 19/50

48000/48000 [=====] - 1s 20us/step - loss: 0.4738 - acc: 0.8343 - val\_loss: 0.4849 - val\_acc: 0.8333

Epoch 20/50

48000/48000 [=====] - 1s 21us/step - loss: 0.4691 - acc: 0.8364 - val\_loss: 0.4831 - val\_acc: 0.8330

Epoch 21/50

48000/48000 [=====] - 1s 21us/step - loss: 0.4655 - acc: 0.8373 - val\_loss: 0.4781 - val\_acc: 0.8349

Epoch 22/50

48000/48000 [=====] - 1s 20us/step - loss: 0.4625 - acc: 0.8374 - val\_loss: 0.4790 - val\_acc: 0.8344

Epoch 23/50

48000/48000 [=====] - 1s 21us/step - loss: 0.4598 - acc: 0.8400 - val\_loss: 0.4719 - val\_acc: 0.8370

Epoch 24/50

```
48000/48000 [=====] - 1s 21us/step - loss: 0.4563 - acc: 0.8400 - val_loss: 0.4691 - val_acc: 0.8380
Epoch 25/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4527 - acc: 0.8421 - val_loss: 0.4739 - val_acc: 0.8372
Epoch 26/50
48000/48000 [=====] - 1s 22us/step - loss: 0.4503 - acc: 0.8420 - val_loss: 0.4665 - val_acc: 0.8373
Epoch 27/50
48000/48000 [=====] - 1s 23us/step - loss: 0.4474 - acc: 0.8436 - val_loss: 0.4620 - val_acc: 0.8401
Epoch 28/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4455 - acc: 0.8439 - val_loss: 0.4609 - val_acc: 0.8403
Epoch 29/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4426 - acc: 0.8461 - val_loss: 0.4661 - val_acc: 0.8357
Epoch 30/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4408 - acc: 0.8459 - val_loss: 0.4592 - val_acc: 0.8416
Epoch 31/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4391 - acc: 0.8474 - val_loss: 0.4551 - val_acc: 0.8418
Epoch 32/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4370 - acc: 0.8478 - val_loss: 0.4527 - val_acc: 0.8430
Epoch 33/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4345 - acc: 0.8482 - val_loss: 0.4537 - val_acc: 0.8420
Epoch 34/50
48000/48000 [=====] - 1s 20us/step - loss: 0.4323 - acc: 0.8491 - val_loss: 0.4547 - val_acc: 0.8409
Epoch 35/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4304 - acc: 0.8500 - val_loss: 0.4514 - val_acc: 0.8432
Epoch 36/50
48000/48000 [=====] - 1s 20us/step - loss: 0.4287 - acc: 0.8494 - val_loss: 0.4450 - val_acc: 0.8464
Epoch 37/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4274 - acc: 0.8499 - val_loss: 0.4478 - val_acc: 0.8425
Epoch 38/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4259 - acc: 0.8510 - val_loss: 0.4436 - val_acc: 0.8456
Epoch 39/50
48000/48000 [=====] - 1s 22us/step - loss: 0.4242 - acc: 0.8515 - val_loss: 0.4447 - val_acc: 0.8445
Epoch 40/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4224 - acc: 0.8523 - val_loss: 0.4516 - val_acc: 0.8419
Epoch 41/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4212 - acc: 0.8522 - val_loss: 0.4557 - val_acc: 0.8372
Epoch 42/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4197 - acc: 0.8530 - val_loss: 0.4459 - val_acc: 0.8455
Epoch 43/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4185 - acc: 0.8541 - val_loss: 0.4445 - val_acc: 0.8437
Epoch 44/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4170 - acc: 0.8539 - val_loss: 0.4406 - val_acc: 0.8453
Epoch 45/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4158 - acc: 0.8548 - val_loss: 0.4342 - val_acc: 0.8482
Epoch 46/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4146 - acc: 0.8538 - val_loss: 0.4365 - val_acc: 0.8484
Epoch 47/50
48000/48000 [=====] - 1s 20us/step - loss: 0.4129 - acc: 0.8551 - val_loss: 0.4342 - val_acc: 0.8472
Epoch 48/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4112 - acc: 0.8565 - val_loss: 0.4324 - val_acc: 0.8472
Epoch 49/50
48000/48000 [=====] - 1s 21us/step - loss: 0.4110 - acc: 0.8548 - val_loss:
```

s: 0.4358 - val\_acc: 0.8477

Epoch 50/50

48000/48000 [=====] - 1s 21us/step - loss: 0.4101 - acc: 0.8566 - val\_loss: 0.4328 - val\_acc: 0.8489

```
In [15]: score = model.evaluate(X_test_mlp, Y_test_mlp, verbose = 1)
print()
print('MLP Test loss:', score[0])
print('MLP Test accuracy:', score[1])
```

10000/10000 [=====] - 0s 50us/step

MLP Test loss: 0.4536403777599335

MLP Test accuracy: 0.8374

## Deep Multi Layer Perceptron

```
In [0]: from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout
from keras.optimizers import RMSprop
from keras.datasets import mnist
from keras.utils import np_utils
```



```
In [17]: # Deep Multilayer Perceptron model
model_deepmlp = Sequential()

model_deepmlp.add(Dense(input_dim=784, units=625, kernel_initializer='normal'))
model_deepmlp.add(Activation('relu'))
model_deepmlp.add(Dropout(0.2))

model_deepmlp.add(Dense(input_dim=625, units=625, kernel_initializer='normal'))
model_deepmlp.add(Activation('relu'))
model_deepmlp.add(Dropout(0.2))

model_deepmlp.add(Dense(input_dim=625, units=625, kernel_initializer='normal'))
model_deepmlp.add(Activation('relu'))
model_deepmlp.add(Dropout(0.2))

model_deepmlp.add(Dense(input_dim=625, units=10, kernel_initializer='normal'))
model_deepmlp.add(Activation('softmax'))

model_deepmlp.compile(optimizer=RMSprop(lr=0.001, rho=0.9),
                      loss='categorical_crossentropy', metrics=['accuracy'])

model_deepmlp.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3445: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 625)	490625
activation_1 (Activation)	(None, 625)	0
dropout_1 (Dropout)	(None, 625)	0
dense_4 (Dense)	(None, 625)	391250
activation_2 (Activation)	(None, 625)	0
dropout_2 (Dropout)	(None, 625)	0
dense_5 (Dense)	(None, 625)	391250
activation_3 (Activation)	(None, 625)	0
dropout_3 (Dropout)	(None, 625)	0
dense_6 (Dense)	(None, 10)	6260
activation_4 (Activation)	(None, 10)	0
Total params: 1,279,385		
Trainable params: 1,279,385		
Non-trainable params: 0		

```
In [18]: history_deepmlp = model_deepmlp.fit(X_train_mlp, Y_train_mlp,  
                                             batch_size = batch_size,  
                                             epochs = num_epochs,  
                                             verbose = 1,  
                                             validation_data = (X_val_mlp, Y_val_mlp))
```

Train on 48000 samples, validate on 12000 samples

Epoch 1/50  
48000/48000 [=====] - 2s 48us/step - loss: 0.6924 - acc: 0.7516 - val\_loss: 0.4532 - val\_acc: 0.8327

Epoch 2/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.4491 - acc: 0.8334 - val\_loss: 0.4684 - val\_acc: 0.8244

Epoch 3/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.4029 - acc: 0.8515 - val\_loss: 0.3793 - val\_acc: 0.8577

Epoch 4/50  
48000/48000 [=====] - 2s 42us/step - loss: 0.3701 - acc: 0.8630 - val\_loss: 0.3968 - val\_acc: 0.8519

Epoch 5/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.3472 - acc: 0.8709 - val\_loss: 0.3730 - val\_acc: 0.8653

Epoch 6/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.3308 - acc: 0.8777 - val\_loss: 0.4574 - val\_acc: 0.8343

Epoch 7/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.3244 - acc: 0.8816 - val\_loss: 0.3770 - val\_acc: 0.8657

Epoch 8/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.3103 - acc: 0.8848 - val\_loss: 0.3821 - val\_acc: 0.8638

Epoch 9/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.3030 - acc: 0.8879 - val\_loss: 0.3970 - val\_acc: 0.8518

Epoch 10/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2967 - acc: 0.8897 - val\_loss: 0.3914 - val\_acc: 0.8653

Epoch 11/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2907 - acc: 0.8944 - val\_loss: 0.4039 - val\_acc: 0.8596

Epoch 12/50  
48000/48000 [=====] - 2s 41us/step - loss: 0.2867 - acc: 0.8943 - val\_loss: 0.3828 - val\_acc: 0.8724

Epoch 13/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2835 - acc: 0.8964 - val\_loss: 0.3726 - val\_acc: 0.8772

Epoch 14/50  
48000/48000 [=====] - 2s 41us/step - loss: 0.2770 - acc: 0.8993 - val\_loss: 0.3600 - val\_acc: 0.8855

Epoch 15/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2713 - acc: 0.9010 - val\_loss: 0.3859 - val\_acc: 0.8720

Epoch 16/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2693 - acc: 0.9032 - val\_loss: 0.4364 - val\_acc: 0.8602

Epoch 17/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2723 - acc: 0.9011 - val\_loss: 0.3938 - val\_acc: 0.8866

Epoch 18/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2666 - acc: 0.9038 - val\_loss: 0.3551 - val\_acc: 0.8853

Epoch 19/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2612 - acc: 0.9052 - val\_loss: 0.3677 - val\_acc: 0.8811

Epoch 20/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2630 - acc: 0.9059 - val\_loss: 0.3541 - val\_acc: 0.8827

Epoch 21/50  
48000/48000 [=====] - 2s 41us/step - loss: 0.2591 - acc: 0.9059 - val\_loss: 0.3798 - val\_acc: 0.8892

Epoch 22/50  
48000/48000 [=====] - 2s 41us/step - loss: 0.2542 - acc: 0.9084 - val\_loss: 0.3815 - val\_acc: 0.8923

Epoch 23/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2539 - acc: 0.9084 - val\_loss: 0.4108 - val\_acc: 0.8827

Epoch 24/50  
48000/48000 [=====] - 2s 41us/step - loss: 0.2562 - acc: 0.9083 - val\_loss: 0.4041 - val\_acc: 0.8888

Epoch 25/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2510 - acc: 0.9101 - val\_loss: 0.3880 - val\_acc: 0.8908

Epoch 26/50  
48000/48000 [=====] - 2s 43us/step - loss: 0.2469 - acc: 0.9120 - val\_loss: 0.4053 - val\_acc: 0.8815  
Epoch 27/50  
48000/48000 [=====] - 2s 42us/step - loss: 0.2484 - acc: 0.9106 - val\_loss: 0.3897 - val\_acc: 0.8869  
Epoch 28/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2410 - acc: 0.9148 - val\_loss: 0.4094 - val\_acc: 0.8841  
Epoch 29/50  
48000/48000 [=====] - 2s 39us/step - loss: 0.2445 - acc: 0.9141 - val\_loss: 0.3949 - val\_acc: 0.8767  
Epoch 30/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2443 - acc: 0.9137 - val\_loss: 0.4137 - val\_acc: 0.8869  
Epoch 31/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2397 - acc: 0.9164 - val\_loss: 0.3946 - val\_acc: 0.8863  
Epoch 32/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2423 - acc: 0.9174 - val\_loss: 0.3852 - val\_acc: 0.8914  
Epoch 33/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2380 - acc: 0.9163 - val\_loss: 0.3963 - val\_acc: 0.8973  
Epoch 34/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2317 - acc: 0.9176 - val\_loss: 0.3979 - val\_acc: 0.8960  
Epoch 35/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2418 - acc: 0.9154 - val\_loss: 0.4151 - val\_acc: 0.8902  
Epoch 36/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2335 - acc: 0.9181 - val\_loss: 0.3894 - val\_acc: 0.8967  
Epoch 37/50  
48000/48000 [=====] - 2s 39us/step - loss: 0.2310 - acc: 0.9186 - val\_loss: 0.4532 - val\_acc: 0.8892  
Epoch 38/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2291 - acc: 0.9191 - val\_loss: 0.4253 - val\_acc: 0.8852  
Epoch 39/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2278 - acc: 0.9200 - val\_loss: 0.4500 - val\_acc: 0.8894  
Epoch 40/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2270 - acc: 0.9213 - val\_loss: 0.4191 - val\_acc: 0.8964  
Epoch 41/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2371 - acc: 0.9199 - val\_loss: 0.5268 - val\_acc: 0.8788  
Epoch 42/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2301 - acc: 0.9207 - val\_loss: 0.4266 - val\_acc: 0.8864  
Epoch 43/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2269 - acc: 0.9228 - val\_loss: 0.3987 - val\_acc: 0.8958  
Epoch 44/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2246 - acc: 0.9228 - val\_loss: 0.3923 - val\_acc: 0.8974  
Epoch 45/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2221 - acc: 0.9238 - val\_loss: 0.4584 - val\_acc: 0.8782  
Epoch 46/50  
48000/48000 [=====] - 2s 39us/step - loss: 0.2225 - acc: 0.9248 - val\_loss: 0.4604 - val\_acc: 0.8788  
Epoch 47/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2256 - acc: 0.9239 - val\_loss: 0.4541 - val\_acc: 0.8908  
Epoch 48/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2204 - acc: 0.9253 - val\_loss: 0.4905 - val\_acc: 0.8909  
Epoch 49/50  
48000/48000 [=====] - 2s 40us/step - loss: 0.2180 - acc: 0.9265 - val\_loss: 0.4492 - val\_acc: 0.8821  
Epoch 50/50  
48000/48000 [=====] - 2s 39us/step - loss: 0.2190 - acc: 0.9262 - val\_loss: 0.4613 - val\_acc: 0.8936

```
In [19]: score = model_deepmlp.evaluate(X_test_mlp, Y_test_mlp, verbose = 1)
print()
print('Deep MLP Test loss:', score[0])
print('Deep MLP Test accuracy:', score[1])

10000/10000 [=====] - 1s 60us/step

Deep MLP Test loss: 0.5017395622313022
Deep MLP Test accuracy: 0.8868
```

## Convolutional Neural Networks

```
In [0]: from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
```

```
In [21]: input_shape = (img_rows, img_cols, 1)

model_cnn = Sequential()
model_cnn.add(Conv2D(32, (3, 3), activation='relu',
                    kernel_initializer='normal', input_shape=input_shape))
model_cnn.add(MaxPooling2D((2, 2)))
model_cnn.add(Dropout(0.25))

model_cnn.add(Conv2D(64, (3, 3), activation='relu'))
model_cnn.add(MaxPooling2D((2, 2)))
model_cnn.add(Dropout(0.25))

model_cnn.add(Conv2D(128, (3, 3), activation='relu'))
model_cnn.add(Dropout(0.4))

model_cnn.add(Flatten())

model_cnn.add(Dense(128, activation='relu'))
model_cnn.add(Dropout(0.3))

model_cnn.add(Dense(num_classes, activation='softmax'))

model_cnn.compile(loss=keras.losses.categorical_crossentropy,
                  optimizer=keras.optimizers.Adam(),
                  metrics=['accuracy'])

model_cnn.summary()
```

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
dropout_4 (Dropout)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
dropout_5 (Dropout)	(None, 5, 5, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 128)	73856
dropout_6 (Dropout)	(None, 3, 3, 128)	0
flatten_1 (Flatten)	(None, 1152)	0
dense_7 (Dense)	(None, 128)	147584
dropout_7 (Dropout)	(None, 128)	0
dense_8 (Dense)	(None, 10)	1290
=====		
Total params: 241,546		
Trainable params: 241,546		
Non-trainable params: 0		

```
In [22]: history_cnn = model_cnn.fit(X_train_cnn, Y_train_cnn,  
                                     batch_size = batch_size,  
                                     epochs = num_epochs,  
                                     verbose = 1,  
                                     validation_data = (X_val_cnn, Y_val_cnn))
```

Train on 48000 samples, validate on 12000 samples

Epoch 1/50  
48000/48000 [=====] - 7s 139us/step - loss: 0.8736 - acc: 0.6707 - val\_loss: 0.5498 - val\_acc: 0.7913

Epoch 2/50  
48000/48000 [=====] - 4s 80us/step - loss: 0.5542 - acc: 0.7915 - val\_loss: 0.4488 - val\_acc: 0.8332

Epoch 3/50  
48000/48000 [=====] - 4s 80us/step - loss: 0.4777 - acc: 0.8215 - val\_loss: 0.3928 - val\_acc: 0.8549

Epoch 4/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.4367 - acc: 0.8394 - val\_loss: 0.3685 - val\_acc: 0.8675

Epoch 5/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.4037 - acc: 0.8513 - val\_loss: 0.3436 - val\_acc: 0.8734

Epoch 6/50  
48000/48000 [=====] - 4s 77us/step - loss: 0.3791 - acc: 0.8603 - val\_loss: 0.3188 - val\_acc: 0.8816

Epoch 7/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.3639 - acc: 0.8670 - val\_loss: 0.3067 - val\_acc: 0.8857

Epoch 8/50  
48000/48000 [=====] - 4s 77us/step - loss: 0.3471 - acc: 0.8725 - val\_loss: 0.3077 - val\_acc: 0.8867

Epoch 9/50  
48000/48000 [=====] - 4s 79us/step - loss: 0.3370 - acc: 0.8747 - val\_loss: 0.2952 - val\_acc: 0.8902

Epoch 10/50  
48000/48000 [=====] - 4s 79us/step - loss: 0.3251 - acc: 0.8800 - val\_loss: 0.2897 - val\_acc: 0.8923

Epoch 11/50  
48000/48000 [=====] - 4s 77us/step - loss: 0.3142 - acc: 0.8839 - val\_loss: 0.2761 - val\_acc: 0.8978

Epoch 12/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.3043 - acc: 0.8877 - val\_loss: 0.2727 - val\_acc: 0.8976

Epoch 13/50  
48000/48000 [=====] - 4s 79us/step - loss: 0.2984 - acc: 0.8913 - val\_loss: 0.2662 - val\_acc: 0.9023

Epoch 14/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2914 - acc: 0.8924 - val\_loss: 0.2656 - val\_acc: 0.9022

Epoch 15/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2838 - acc: 0.8968 - val\_loss: 0.2677 - val\_acc: 0.9008

Epoch 16/50  
48000/48000 [=====] - 4s 77us/step - loss: 0.2797 - acc: 0.8970 - val\_loss: 0.2557 - val\_acc: 0.9060

Epoch 17/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2743 - acc: 0.8977 - val\_loss: 0.2546 - val\_acc: 0.9070

Epoch 18/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2691 - acc: 0.9016 - val\_loss: 0.2662 - val\_acc: 0.8978

Epoch 19/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2643 - acc: 0.9011 - val\_loss: 0.2558 - val\_acc: 0.9074

Epoch 20/50  
48000/48000 [=====] - 4s 79us/step - loss: 0.2604 - acc: 0.9032 - val\_loss: 0.2401 - val\_acc: 0.9116

Epoch 21/50  
48000/48000 [=====] - 4s 77us/step - loss: 0.2544 - acc: 0.9052 - val\_loss: 0.2380 - val\_acc: 0.9135

Epoch 22/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2525 - acc: 0.9066 - val\_loss: 0.2366 - val\_acc: 0.9117

Epoch 23/50  
48000/48000 [=====] - 4s 77us/step - loss: 0.2485 - acc: 0.9082 - val\_loss: 0.2394 - val\_acc: 0.9123

Epoch 24/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2490 - acc: 0.9072 - val\_loss: 0.2378 - val\_acc: 0.9122

Epoch 25/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2424 - acc: 0.9086 - val\_loss: 0.2274 - val\_acc: 0.9153



Epoch 26/50  
48000/48000 [=====] - 4s 77us/step - loss: 0.2389 - acc: 0.9101 - val\_loss: 0.2333 - val\_acc: 0.9140  
Epoch 27/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2362 - acc: 0.9126 - val\_loss: 0.2377 - val\_acc: 0.9126  
Epoch 28/50  
48000/48000 [=====] - 4s 76us/step - loss: 0.2311 - acc: 0.9140 - val\_loss: 0.2265 - val\_acc: 0.9141  
Epoch 29/50  
48000/48000 [=====] - 4s 79us/step - loss: 0.2328 - acc: 0.9142 - val\_loss: 0.2307 - val\_acc: 0.9124  
Epoch 30/50  
48000/48000 [=====] - 4s 79us/step - loss: 0.2296 - acc: 0.9136 - val\_loss: 0.2292 - val\_acc: 0.9158  
Epoch 31/50  
48000/48000 [=====] - 4s 77us/step - loss: 0.2250 - acc: 0.9165 - val\_loss: 0.2303 - val\_acc: 0.9147  
Epoch 32/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2227 - acc: 0.9162 - val\_loss: 0.2233 - val\_acc: 0.9162  
Epoch 33/50  
48000/48000 [=====] - 4s 77us/step - loss: 0.2216 - acc: 0.9180 - val\_loss: 0.2486 - val\_acc: 0.9069  
Epoch 34/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2186 - acc: 0.9181 - val\_loss: 0.2213 - val\_acc: 0.9177  
Epoch 35/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2151 - acc: 0.9196 - val\_loss: 0.2278 - val\_acc: 0.9147  
Epoch 36/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2176 - acc: 0.9191 - val\_loss: 0.2331 - val\_acc: 0.9138  
Epoch 37/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2130 - acc: 0.9199 - val\_loss: 0.2203 - val\_acc: 0.9186  
Epoch 38/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2130 - acc: 0.9199 - val\_loss: 0.2208 - val\_acc: 0.9201  
Epoch 39/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2071 - acc: 0.9202 - val\_loss: 0.2239 - val\_acc: 0.9174  
Epoch 40/50  
48000/48000 [=====] - 4s 79us/step - loss: 0.2055 - acc: 0.9218 - val\_loss: 0.2197 - val\_acc: 0.9178  
Epoch 41/50  
48000/48000 [=====] - 4s 79us/step - loss: 0.2050 - acc: 0.9237 - val\_loss: 0.2245 - val\_acc: 0.9189  
Epoch 42/50  
48000/48000 [=====] - 4s 80us/step - loss: 0.2049 - acc: 0.9238 - val\_loss: 0.2279 - val\_acc: 0.9162  
Epoch 43/50  
48000/48000 [=====] - 4s 77us/step - loss: 0.2005 - acc: 0.9247 - val\_loss: 0.2197 - val\_acc: 0.9193  
Epoch 44/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.2016 - acc: 0.9233 - val\_loss: 0.2251 - val\_acc: 0.9173  
Epoch 45/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.1986 - acc: 0.9245 - val\_loss: 0.2185 - val\_acc: 0.9213  
Epoch 46/50  
48000/48000 [=====] - 4s 77us/step - loss: 0.1965 - acc: 0.9246 - val\_loss: 0.2208 - val\_acc: 0.9173  
Epoch 47/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.1988 - acc: 0.9251 - val\_loss: 0.2224 - val\_acc: 0.9191  
Epoch 48/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.1965 - acc: 0.9265 - val\_loss: 0.2183 - val\_acc: 0.9193  
Epoch 49/50  
48000/48000 [=====] - 4s 78us/step - loss: 0.1931 - acc: 0.9287 - val\_loss: 0.2150 - val\_acc: 0.9217  
Epoch 50/50  
48000/48000 [=====] - 4s 79us/step - loss: 0.1913 - acc: 0.9279 - val\_loss: 0.2177 - val\_acc: 0.9214

```
In [23]: score = model_cnn.evaluate(X_test_cnn, Y_test_cnn, verbose = 1)
print()
print('CNN Test loss:', score[0])
print('CNN Test accuracy:', score[1])

10000/10000 [=====] - 1s 82us/step

CNN Test loss: 0.22900302994847296
CNN Test accuracy: 0.9187
```


## Results

```
In [24]: accuracy = history_cnn.history['acc']
val_accuracy = history_cnn.history['val_acc']

loss = history_cnn.history['loss']
val_loss = history_cnn.history['val_loss']

epochs = range(len(accuracy))
plt.plot(epochs, accuracy, 'b', label='Training accuracy')
plt.plot(epochs, val_accuracy, 'r', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()
plt.plot(epochs, loss, 'b', label='Training loss')
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



## Classification Report

```
In [25]: # Get the predictions for the test data
predicted_classes = model_cnn.predict_classes(X_test_cnn)

# Get the indices to be plotted
Y_true_cnn = Y_test
correct = np.nonzero(predicted_classes == Y_true_cnn)
incorrect = np.nonzero(predicted_classes != Y_true_cnn)

from sklearn.metrics import classification_report

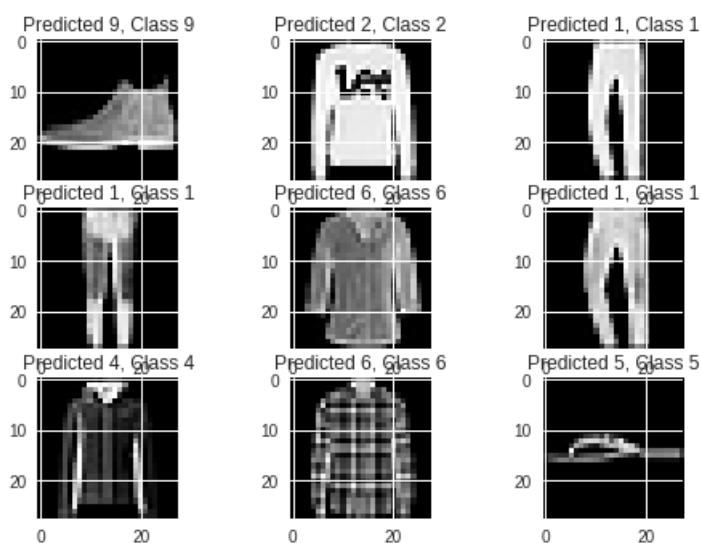
target_names = ["Class {}".format(i) for i in range(num_classes)]

print(classification_report(Y_true_cnn, predicted_classes, target_names = target_names))
```

	precision	recall	f1-score	support
Class 0	0.85	0.89	0.87	1000
Class 1	0.99	0.98	0.99	1000
Class 2	0.85	0.91	0.88	1000
Class 3	0.91	0.93	0.92	1000
Class 4	0.89	0.88	0.88	1000
Class 5	0.99	0.98	0.98	1000
Class 6	0.80	0.70	0.75	1000
Class 7	0.94	0.99	0.96	1000
Class 8	0.97	0.98	0.98	1000
Class 9	0.99	0.95	0.97	1000
micro avg	0.92	0.92	0.92	10000
macro avg	0.92	0.92	0.92	10000
weighted avg	0.92	0.92	0.92	10000

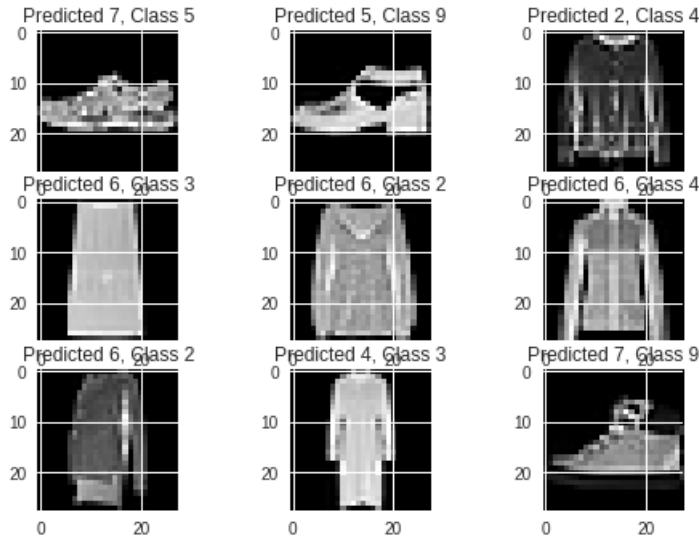
## Subset of correctly predicted classes:

```
In [26]: i=1
for j in range(9):
    plt.subplot(3,3,i)
    plt.imshow(X_test_cnn[correct[0][j]].reshape(28,28), cmap='gray', interpolation='none')
    plt.title("Predicted {}, Class {}".format(predicted_classes[correct[0][j]], Y_true_cnn[correct[0][j]]))
    i+=1
```



## Subset of incorrectly predicted classes:

```
In [27]: i=1
for j in range(9):
    plt.subplot(3,3,i)
    plt.imshow(X_test_cnn[incorrect[0][j]].reshape(28,28), cmap='gray', interpolation='none')
    plt.title("Predicted {}, Class {}".format(predicted_classes[incorrect[0][j]], Y_true_cnn[incorrect[0][j]]))
    i+=1
```



It looks like diversity of the similar patterns present on multiple classes effect the performance of the classifier although CNN is a robust architecture. A jacket, a shirt, and a long-sleeve blouse has similar patterns: long sleeves (or not!), buttons (or not!), and so on.

```
In [0]: # Confusion Matrix
from sklearn.metrics import confusion_matrix
import itertools

def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```


```
In [29]: # Predict the values from the validation dataset
Y_pred = model_cnn.predict(X_test_cnn)

# Convert predictions classes to one hot vectors
Y_pred_classes = np.argmax(Y_pred,axis = 1)

# Convert validation observations to one hot vectors
Y_true = np.argmax(Y_test_cnn,axis = 1)

# compute the confusion matrix
confusion_mtx = confusion_matrix(Y_true, Y_pred_classes)

# plot the confusion matrix
plot_confusion_matrix(confusion_mtx, classes = range(10))
```



## Transfer Learning

Based on our past experience, we humans can learn a new skill easily. We are more efficient in learning, particularly if the task in hand similar to what we have done have done in similar, for example learning a new programming language for a computer professional or driving a new type of vehicle for a seasoned driver is relatively easy based on our past experience.

Transfer learning is an area in machine learning that aims to utilize the knowledge gained while solving one problem to solve a different but related problem.

We will train a CNN model of two level layers i.e., a **feature layer** and a **classification layer** on the first 5 digits (0 to 4) of MNIST dataset

Apply **Transfer Learning** to **freeze features layer** and **fine-tune dense layers** for the classification of digits 5 to 9.

```
In [30]: (X_train, Y_train), (X_test, Y_test) = tf.keras.datasets.mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11493376/11490434 [=====] - 0s 0us/step
```

```
In [31]: print(X_train.shape)
print(Y_train.shape)
print(X_test.shape)
print(Y_test.shape)
```

```
(60000, 28, 28)
(60000,)
(10000, 28, 28)
(10000,)
```

In [0]: *# Create two datasets: one with digits below 5 and one with 5 and above*

```
X_train_lt5 = X_train[Y_train < 5]
Y_train_lt5 = Y_train[Y_train < 5]
X_test_lt5 = X_test[Y_test < 5]
Y_test_lt5 = Y_test[Y_test < 5]

X_train_gte5 = X_train[Y_train >= 5]
Y_train_gte5 = Y_train[Y_train >= 5] - 5 # make classes start at 0
X_test_gte5 = X_test[Y_test >= 5]
Y_test_gte5 = Y_test[Y_test >= 5] - 5
```

In [0]: *# Initialize variables*

```
batch_size = 128
num_classes = 5 # for training stage of transfer learning
epochs = 5

# number of convolutional filters to use
num_filters = 32

# size of pooling area for max pooling
pool_size = 2

# convolution kernel size
kernel_size = 3
```

### Train model for digits 0 to 4

In [0]: *# Function for training model for transfer learning*

```
def train_model(model, train, test, num_classes):
    # reshape
    X_train = train[0].reshape((train[0].shape[0],) + input_shape)
    X_test = test[0].reshape((test[0].shape[0],) + input_shape)

    # convert data type
    X_train = X_train.astype('float32')
    X_test = X_test.astype('float32')

    # normalize pixel values
    X_train /= 255
    X_test /= 255

    # convert class vectors to binary class matrices
    Y_train = np_utils.to_categorical(train[1], num_classes)
    Y_test = np_utils.to_categorical(test[1], num_classes)

    # compile
    model.compile(loss='categorical_crossentropy',
                  optimizer='adadelta',
                  metrics=['accuracy'])

    # fitting of model
    model.fit(X_train, Y_train,
              batch_size=batch_size, epochs=epochs,
              verbose=1,
              validation_data=(X_test, Y_test))

    # evaluation
    score = model.evaluate(X_test, Y_test, verbose=1)

    # report scores
    print()
    print('Test score:', score[0])
    print('Test accuracy:', score[1])
    print()
```

```
In [35]: from keras.layers import Convolution2D, MaxPooling2D
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten

# Define two groups of layers: feature (convolutions) and classification (dense)
feature_layers = [
    Convolution2D(num_filters, (kernel_size, kernel_size),
                  input_shape = input_shape,
                  padding = 'valid'),
    Activation('relu'),
    Convolution2D(num_filters, (kernel_size, kernel_size)),
    Activation('relu'),
    MaxPooling2D((pool_size, pool_size)),
    Dropout(0.25),
    Flatten(),
]

classification_layers = [
    Dense(128),
    Activation('relu'),
    Dropout(0.5),
    Dense(num_classes),
    Activation('softmax')
]

# Create complete model
model = Sequential(feature_layers + classification_layers)
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 26, 26, 32)	320
activation_5 (Activation)	(None, 26, 26, 32)	0
conv2d_5 (Conv2D)	(None, 24, 24, 32)	9248
activation_6 (Activation)	(None, 24, 24, 32)	0
max_pooling2d_3 (MaxPooling2D)	(None, 12, 12, 32)	0
dropout_8 (Dropout)	(None, 12, 12, 32)	0
flatten_2 (Flatten)	(None, 4608)	0
dense_9 (Dense)	(None, 128)	589952
activation_7 (Activation)	(None, 128)	0
dropout_9 (Dropout)	(None, 128)	0
dense_10 (Dense)	(None, 5)	645
activation_8 (Activation)	(None, 5)	0
Total params: 600,165		
Trainable params: 600,165		
Non-trainable params: 0		

```
In [36]: # Train model for 5-digit classification [0 to 4]
train_model(model, (X_train_lt5, Y_train_lt5), (X_test_lt5, Y_test_lt5), num_classes)

Train on 30596 samples, validate on 5139 samples
Epoch 1/5
30596/30596 [=====] - 4s 133us/step - loss: 0.1532 - acc: 0.9510 - val_loss: 0.0252 - val_acc: 0.9914
Epoch 2/5
30596/30596 [=====] - 3s 113us/step - loss: 0.0433 - acc: 0.9869 - val_loss: 0.0136 - val_acc: 0.9955
Epoch 3/5
30596/30596 [=====] - 3s 110us/step - loss: 0.0296 - acc: 0.9910 - val_loss: 0.0119 - val_acc: 0.9957
Epoch 4/5
30596/30596 [=====] - 3s 111us/step - loss: 0.0261 - acc: 0.9923 - val_loss: 0.0084 - val_acc: 0.9973
Epoch 5/5
30596/30596 [=====] - 3s 110us/step - loss: 0.0204 - acc: 0.9940 - val_loss: 0.0082 - val_acc: 0.9967
5139/5139 [=====] - 0s 81us/step

Test score: 0.008182254664061858
Test accuracy: 0.9966919634170072
```

### Transfer existing trained model on 0 to 4 to build model for digits 5 to 9

```
In [37]: # Freeze feature layers and rebuild model
for layer in feature_layers:
    layer.trainable = False

# Transfer: train dense layers for new classification task [5 to 9]
train_model(model, (X_train_gte5, Y_train_gte5), (X_test_gte5, Y_test_gte5), num_classes)

Train on 29404 samples, validate on 4861 samples
Epoch 1/5
29404/29404 [=====] - 3s 90us/step - loss: 0.2423 - acc: 0.9300 - val_loss: 0.0507 - val_acc: 0.9844
Epoch 2/5
29404/29404 [=====] - 2s 68us/step - loss: 0.0721 - acc: 0.9778 - val_loss: 0.0350 - val_acc: 0.9889
Epoch 3/5
29404/29404 [=====] - 2s 68us/step - loss: 0.0561 - acc: 0.9828 - val_loss: 0.0295 - val_acc: 0.9897
Epoch 4/5
29404/29404 [=====] - 2s 67us/step - loss: 0.0465 - acc: 0.9855 - val_loss: 0.0252 - val_acc: 0.9907
Epoch 5/5
29404/29404 [=====] - 2s 69us/step - loss: 0.0408 - acc: 0.9877 - val_loss: 0.0252 - val_acc: 0.9916
4861/4861 [=====] - 0s 88us/step

Test score: 0.025249082497051406
Test accuracy: 0.9915655214976342
```