

Python Libraries for ML @ GCE Raipur

Instructor: Santosh Chapaneri

Mar 2021

Scikit-learn Library for ML

Logistic Regression

```
In [1]: import numpy as np
import pandas as pd
import sklearn
import sklearn.linear_model as lm
import sklearn.model_selection as cv
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [3]: titanic = pd.read_csv('titanic.csv')
titanic.head()
```

Out[3]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

- We will keep only a few fields for this example.
- We also convert the sex field to a binary variable, so that it can be handled correctly by NumPy and scikit-learn.
- Finally, we remove the rows containing NaN values.

```
In [4]: data = titanic[['Sex', 'Age', 'Pclass', 'Survived']].copy()
data['Sex'] = data['Sex'] == 'female'
data = data.dropna()
data.head()
```

Out[4]:

	Sex	Age	Pclass	Survived
0	False	22.0	3	0
1	True	38.0	1	1
2	True	26.0	3	1
3	True	35.0	1	1
4	False	35.0	3	0

- Now, we convert this DataFrame to a NumPy array, so that we can pass it to scikit-learn.

```
In [5]: data_np = data.astype(np.int32).values
X = data_np[:, :-1] # Features
y = data_np[:, -1] # Target (survived or not)
print(X[:5, :])
print(y[:5])

[[ 0 22  3]
 [ 1 38  1]
 [ 1 26  3]
 [ 1 35  1]
 [ 0 35  3]]
[[0 1 1 1 0]]
```

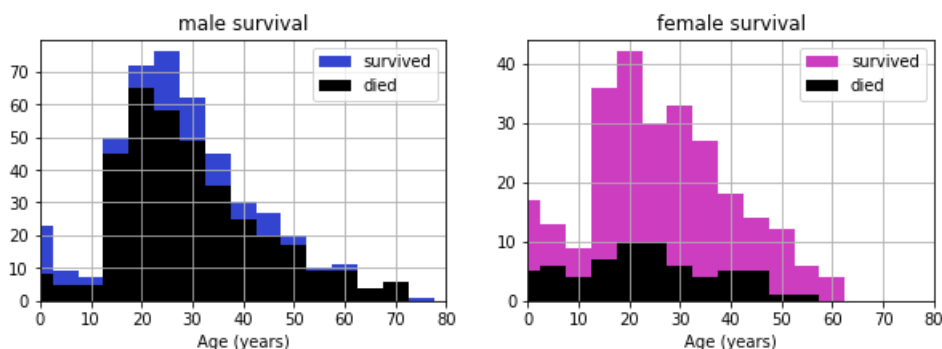
- Let's have a look at the survival of male and female passengers, as a function of their age:

```
In [6]: # Need few boolean vectors
female = (X[:, 0] == 1)
survived = (y == 1)

# This vector contains the age of the passengers.
age = X[:, 1]

# Compute few histograms
mybins = np.arange(0, 81, 5)
S = {'male': np.histogram(age[survived & ~female],
                           bins=mybins)[0],
      'female': np.histogram(age[survived & female],
                              bins=mybins)[0]}
D = {'male': np.histogram(age[~survived & ~female],
                           bins=mybins)[0],
      'female': np.histogram(age[~survived & female],
                              bins=mybins)[0]}
```

```
In [7]: # Now plot the data
bins = mybins[:-1]
plt.figure(figsize=(10, 3));
for i, sex, color in zip((0, 1),
                        ('male', 'female'),
                        ('#3345d0', '#cc3dc0')):
    plt.subplot(121 + i);
    plt.bar(bins, S[sex], bottom=D[sex], color=color,
            width=5, label='survived');
    plt.bar(bins, D[sex], color='k', width=5, label='died');
    plt.xlim(0, 80);
    plt.grid(None);
    plt.title(sex + " survival");
    plt.xlabel("Age (years)");
    plt.legend();
```



- Let's train a LogisticRegression classifier. We first need to create a train and a test dataset.

```
In [8]: # Split X and y into train and test datasets
(X_train, X_test, y_train, y_test) = cv.train_test_split(X, y, test_size=.25)
```

```
In [9]: # Instantiate the classifier / 'Create the model'
logreg = lm.LogisticRegression()
```

- Train the model (fit) and get the predicted values (predict) on the test set.

```
In [10]: logreg.fit(X_train, y_train)
```

```
Out[10]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='auto', n_jobs=None, penalty='l2',
random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
warm_start=False)
```

```
In [11]: y_predicted = logreg.predict(X_test)
```

```
In [12]: from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_predicted)
```

```
Out[12]: 0.8156424581005587
```

```
In [13]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_predicted))
```

	precision	recall	f1-score	support
0	0.84	0.85	0.85	108
1	0.77	0.76	0.77	71
accuracy			0.82	179
macro avg	0.81	0.81	0.81	179
weighted avg	0.82	0.82	0.82	179

```
In [14]: from sklearn.metrics import confusion_matrix
print(confusion_matrix(y_test, y_predicted))
```

```
[[92 16]
 [17 54]]
```

Run Logistic Regression on the IRIS Dataset

```
In [15]: import sklearn.datasets as ds
iris = ds.load_iris()

X, y = iris.data, iris.target
(X_train, X_test, y_train, y_test) = cv.train_test_split(X, y, test_size=.25)

logreg = lm.LogisticRegression()
logreg.fit(X_train, y_train)
y_predicted = logreg.predict(X_test)

print(accuracy_score(y_test, y_predicted))
print(classification_report(y_test, y_predicted))
print(confusion_matrix(y_test, y_predicted))
```

```
1.0

              precision    recall  f1-score   support

     0         1.00        1.00        1.00        15
     1         1.00        1.00        1.00        14
     2         1.00        1.00        1.00         9

 accuracy          1.00
 macro avg          1.00
weighted avg          1.00
```

```
[[15  0  0]
 [ 0 14  0]
 [ 0  0  9]]
```

/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning:
lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

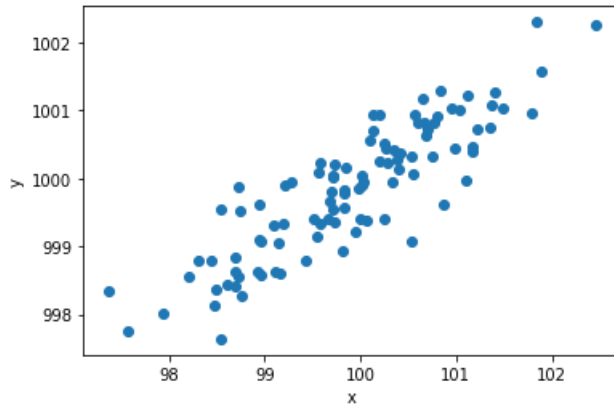
Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

Linear Regression

```
In [16]: import numpy as np
import pandas as pd
import sklearn
import sklearn.linear_model as lm
import sklearn.model_selection as cv
import matplotlib.pyplot as plt
import scipy.stats
%matplotlib inline
```

```
In [17]: rng = np.random.RandomState(123)
mean = [100, 1000]
cov = [[1, 0.9], [0.9, 1]]
sample = rng.multivariate_normal(mean, cov, size=100)
x, y = sample[:, 0], sample[:, 1]

plt.scatter(x, y)
plt.xlabel('x')
plt.ylabel('y')
plt.show()
```



```
In [18]: # x.shape => (100,)
# newaxis: increase the dimension of existing array by one more dimension
X = x[:, np.newaxis]
# X.shape => (100,1)
X.shape
```

Out[18]: (100, 1)

```
In [19]: # adding a column vector of "ones"
# hstack: stack arrays in sequence horizontally (column wise)
Xb = np.hstack((np.ones((X.shape[0], 1)), X))
# Xb.shape => (100, 2); first column for bias, second column for X
w = np.zeros(Xb.shape[1])

# Closed-form solution
z = np.linalg.inv(np.dot(Xb.T, Xb))

w = np.dot(z, np.dot(Xb.T, y))

b, w1 = w[0], w[1]

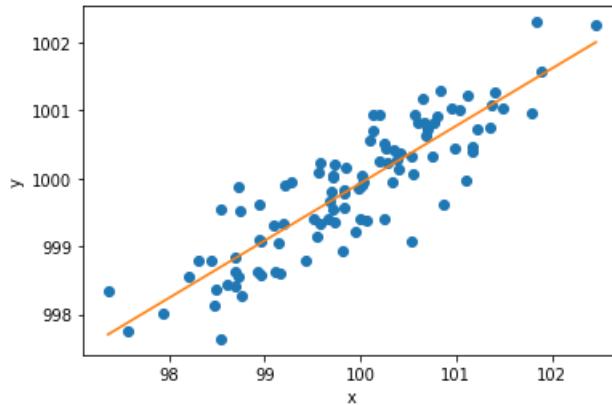
print('slope: %.2f' % w1)
print('y-intercept: %.2f' % b)
```

```
slope: 0.84
y-intercept: 915.59
```

- Show the line fit:

```
In [20]: extremes = np.array([np.min(x), np.max(x)])
predict = extremes*w1 + b

plt.plot(x, y, marker='o', linestyle='')
plt.plot(extremes, predict)
plt.xlabel('x')
plt.ylabel('y')
plt.show()
```



```
In [21]: y_predicted = x*w1 + b

mse = np.mean((y - y_predicted)**2)
mse
```

Out[21]: 0.21920128791623675

Support Vector Machines (SVM)

```
In [22]: import numpy as np
import pandas as pd
import sklearn.datasets as ds
import sklearn.model_selection as cv
import sklearn.svm as svm
import matplotlib as mpl
import matplotlib.pyplot as plt
%matplotlib inline
```

- We generate 2D points and assign a binary label according to a linear operation on the coordinates.

```
In [23]: X = np.random.randn(200, 2)
y = (X[:, 0] + X[:, 1]) > 1 # Imagine an AND function

X[:10], y[:10]
```

```
Out[23]: (array([[ 0.10099719, -0.29817764],
 [-0.25311934,  0.27369246],
 [ 1.44501533, -0.42040294],
 [ 1.21945629, -0.11666772],
 [-1.09791891,  2.50445651],
 [ 0.07068561,  0.75997489],
 [-1.991155   ,  1.06416644],
 [-1.69995175, -1.35561656],
 [ 1.3107086 , -2.1062814 ],
 [-0.12962302, -0.01251449]]),
 array([False, False,  True,  True,  True, False, False, False, False,
        False]))
```

- Fit a linear Support Vector Classifier (SVC)

```
In [24]: est = svm.LinearSVC()
est.fit(X, y)
```

```
Out[24]: LinearSVC(C=1.0, class_weight=None, dual=True, fit_intercept=True,
intercept_scaling=1, loss='squared_hinge', max_iter=1000,
multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
verbose=0)
```

```
In [25]: # Generate a grid in the square  $[-3, 3]^2$ 
xx, yy = np.meshgrid(np.linspace(-3, 3, 500),
                    np.linspace(-3, 3, 500))

# This function takes a SVM estimator as input
def plot_decision_function(est, X, y):
    # We evaluate the decision function on the grid.
    Z = est.decision_function(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    cmap = plt.cm.Blues

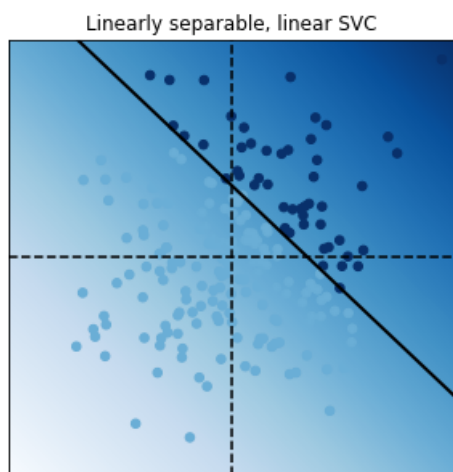
    # Display the decision function on the grid
    plt.figure(figsize=(5,5));
    plt.imshow(Z,
               extent=(xx.min(), xx.max(), yy.min(), yy.max()),
               aspect='auto', origin='lower', cmap=cmap)

    # Display the boundaries
    plt.contour(xx, yy, Z, levels=[0], linewidths=2, colors='k')

    # Display the points with their true labels
    plt.scatter(X[:, 0], X[:, 1], s=30, c=.5+.5*y, lw=1,
               cmap=cmap, vmin=0, vmax=1);
    plt.axhline(0, color='k', ls='--')
    plt.axvline(0, color='k', ls='--')
    plt.xticks(())
    plt.yticks(())
    plt.axis([-3, 3, -3, 3])
```

```
In [26]: plot_decision_function(est, X, y)
plt.title("Linearly separable, linear SVC")
```

```
Out[26]: Text(0.5, 1.0, 'Linearly separable, linear SVC')
```



The linear SVC tried to separate the points with a line and it did a good job.

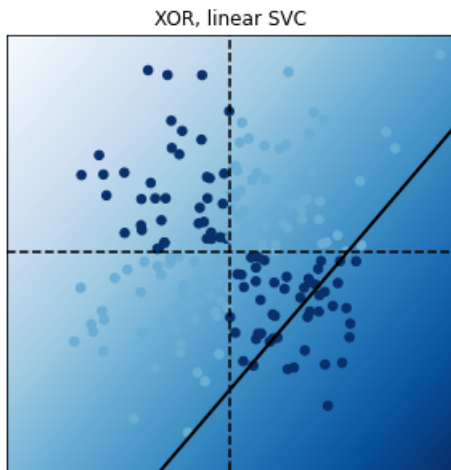
- We now modify the labels with a XOR function.
- A point's label is 1 if the coordinates have different signs. This classification is not linearly separable. Therefore, a linear SVC fails completely.

```
In [27]: y = np.logical_xor(X[:, 0] > 0, X[:, 1] > 0)

est2 = svm.LinearSVC()
est2.fit(X, y)

plot_decision_function(est2, X, y)
plt.title("XOR, linear SVC")
```

Out[27]: Text(0.5, 1.0, 'XOR, linear SVC')



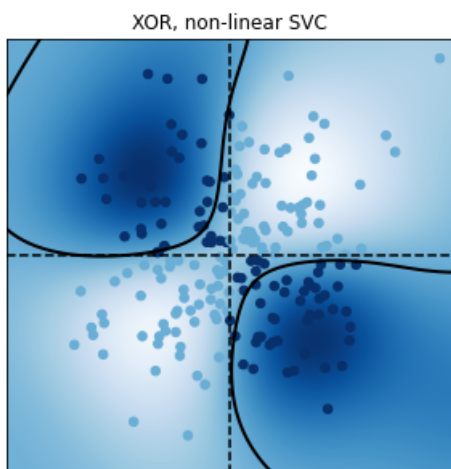
- It is possible to use non-linear SVCs by using non-linear kernels.
- Kernels specify a non-linear transformation of the points into a higher-dimensional space. Transformed points in this space are assumed to be more linearly separable, although they are not necessarily in the original space.
- By default, the SVC classifier in scikit-learn uses the Radial Basis Function (RBF) kernel.

```
In [28]: y = np.logical_xor(X[:, 0] > 0, X[:, 1] > 0)

est3 = svm.SVC()
est3.fit(X, y)

plot_decision_function(est3, X, y)
plt.title("XOR, non-linear SVC")
```

Out[28]: Text(0.5, 1.0, 'XOR, non-linear SVC')



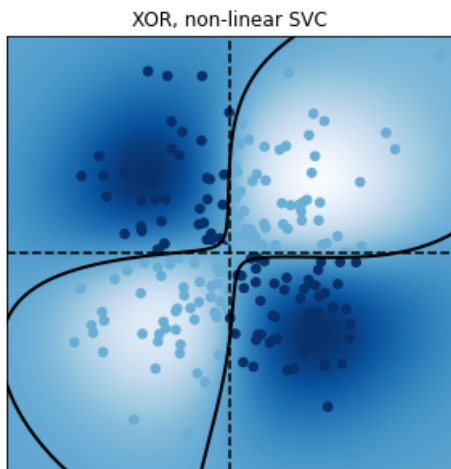

```
In [29]: from sklearn.model_selection import GridSearchCV

est4 = GridSearchCV(svm.SVC(),
                    {'C': np.logspace(-3., 3., 10),
                     'gamma': np.logspace(-3., 3., 10)}, cv=5);
est4.fit(X, y)
print("Score: {0:.3f}".format(
    cv.cross_val_score(est4, X, y).mean()))

plot_decision_function(est4.best_estimator_, X, y)
plt.title("XOR, non-linear SVC")
```

Score: 0.940

Out[29]: Text(0.5, 1.0, 'XOR, non-linear SVC')



Applying SVM on IRIS dataset

```
In [30]: import pandas as pd
import numpy as np
from sklearn import svm, datasets
import matplotlib.pyplot as plt
%matplotlib inline

iris = datasets.load_iris()
X = iris.data
y = iris.target
```

- Let's do **hyper-parameter tuning**
- Use **5-fold cross validation** to perform grid search to calculate optimal hyper-parameters

```
In [31]: from sklearn.model_selection import GridSearchCV
import sklearn.model_selection as cv
from sklearn.metrics import classification_report
from sklearn.utils import shuffle

# Split the dataset
X_train, X_test, y_train, y_test = cv.train_test_split(X, y, test_size=0.25)
```

```
In [32]: # Set the parameters by cross-validation
parameters = [{'kernel': ['rbf'],
                      'gamma': [1e-4, 1e-3, 0.01, 0.1, 0.2, 0.5],
                      'C': [1, 10, 100, 1000]},
               {'kernel': ['linear'], 'C': [1, 10, 100, 1000]}]

print("# Tuning hyper-parameters")
print()

clf = GridSearchCV(svm.SVC(decision_function_shape='ovr'), parameters, cv=5)
clf.fit(X_train, y_train)
```

Tuning hyper-parameters

```
Out[32]: GridSearchCV(cv=5, error_score=nan,
                      estimator=SVC(C=1.0, break_ties=False, cache_size=200,
                                     class_weight=None, coef0=0.0,
                                     decision_function_shape='ovr', degree=3,
                                     gamma='scale', kernel='rbf', max_iter=-1,
                                     probability=False, random_state=None, shrinking=True,
                                     tol=0.001, verbose=False),
                      iid='deprecated', n_jobs=None,
                      param_grid=[{'C': [1, 10, 100, 1000],
                                   'gamma': [0.0001, 0.001, 0.01, 0.1, 0.2, 0.5],
                                   'kernel': ['rbf']},
                                   {'C': [1, 10, 100, 1000], 'kernel': ['linear']}],
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                      scoring=None, verbose=0)
```

```
In [33]: print("Best parameters set found on training set:")
print()
print(clf.best_params_)
print()
print("Grid scores on training set:")
print()
means = clf.cv_results_['mean_test_score']
stds = clf.cv_results_['std_test_score']
for mean, std, params in zip(means, stds, clf.cv_results_['params']):
    print("%0.3f (+/-%0.03f) for %r"
          % (mean, std * 2, params))
print()
```

Best parameters set found on training set:

```
{'C': 1, 'gamma': 0.5, 'kernel': 'rbf'}
```

Grid scores on training set:

```
0.357 (+/-0.015) for {'C': 1, 'gamma': 0.0001, 'kernel': 'rbf'}
0.687 (+/-0.014) for {'C': 1, 'gamma': 0.001, 'kernel': 'rbf'}
0.919 (+/-0.157) for {'C': 1, 'gamma': 0.01, 'kernel': 'rbf'}
0.955 (+/-0.141) for {'C': 1, 'gamma': 0.1, 'kernel': 'rbf'}
0.973 (+/-0.072) for {'C': 1, 'gamma': 0.2, 'kernel': 'rbf'}
0.974 (+/-0.070) for {'C': 1, 'gamma': 0.5, 'kernel': 'rbf'}
0.687 (+/-0.014) for {'C': 10, 'gamma': 0.0001, 'kernel': 'rbf'}
0.919 (+/-0.157) for {'C': 10, 'gamma': 0.001, 'kernel': 'rbf'}
0.955 (+/-0.141) for {'C': 10, 'gamma': 0.01, 'kernel': 'rbf'}
0.964 (+/-0.087) for {'C': 10, 'gamma': 0.1, 'kernel': 'rbf'}
0.955 (+/-0.114) for {'C': 10, 'gamma': 0.2, 'kernel': 'rbf'}
0.955 (+/-0.114) for {'C': 10, 'gamma': 0.5, 'kernel': 'rbf'}
0.919 (+/-0.157) for {'C': 100, 'gamma': 0.0001, 'kernel': 'rbf'}
0.955 (+/-0.141) for {'C': 100, 'gamma': 0.001, 'kernel': 'rbf'}
0.955 (+/-0.114) for {'C': 100, 'gamma': 0.01, 'kernel': 'rbf'}
0.955 (+/-0.114) for {'C': 100, 'gamma': 0.1, 'kernel': 'rbf'}
0.964 (+/-0.087) for {'C': 100, 'gamma': 0.2, 'kernel': 'rbf'}
0.955 (+/-0.114) for {'C': 100, 'gamma': 0.5, 'kernel': 'rbf'}
0.955 (+/-0.141) for {'C': 1000, 'gamma': 0.0001, 'kernel': 'rbf'}
0.955 (+/-0.114) for {'C': 1000, 'gamma': 0.001, 'kernel': 'rbf'}
0.955 (+/-0.114) for {'C': 1000, 'gamma': 0.01, 'kernel': 'rbf'}
0.974 (+/-0.070) for {'C': 1000, 'gamma': 0.1, 'kernel': 'rbf'}
0.956 (+/-0.080) for {'C': 1000, 'gamma': 0.2, 'kernel': 'rbf'}
0.946 (+/-0.105) for {'C': 1000, 'gamma': 0.5, 'kernel': 'rbf'}
0.946 (+/-0.145) for {'C': 1, 'kernel': 'linear'}
0.955 (+/-0.114) for {'C': 10, 'kernel': 'linear'}
0.955 (+/-0.114) for {'C': 100, 'kernel': 'linear'}
0.955 (+/-0.114) for {'C': 1000, 'kernel': 'linear'}
```

```
In [34]: print("Detailed classification report:")
print()
print("The model is trained on the training set.")
print("The scores are computed on the testing set.")
print()
y_true, y_pred = y_test, clf.predict(X_test)
print(classification_report(y_true, y_pred))
print()
# The support is the number of occurrences of each class in y_true
```

Detailed classification report:

The model is trained on the training set.
The scores are computed on the testing set.

	precision	recall	f1-score	support
0	1.00	1.00	1.00	13
1	0.93	0.93	0.93	15
2	0.90	0.90	0.90	10
accuracy			0.95	38
macro avg	0.94	0.94	0.94	38
weighted avg	0.95	0.95	0.95	38

```
In [35]: from sklearn.metrics import accuracy_score
accuracy_score(y_true, y_pred)
```

```
Out[35]: 0.9473684210526315
```

K-Means Clustering

- Here we'll explore K Means Clustering, which is an unsupervised clustering technique.
- K-Means is an algorithm for unsupervised clustering: that is, finding clusters in data based on the data attributes alone (not the labels).
- K-Means is a relatively easy-to-understand algorithm. It searches for cluster centers which are the mean of the points within them, such that every point is closest to the cluster center it is assigned to.

```
In [36]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from scipy import stats

# use seaborn plotting defaults
import seaborn as sns
```

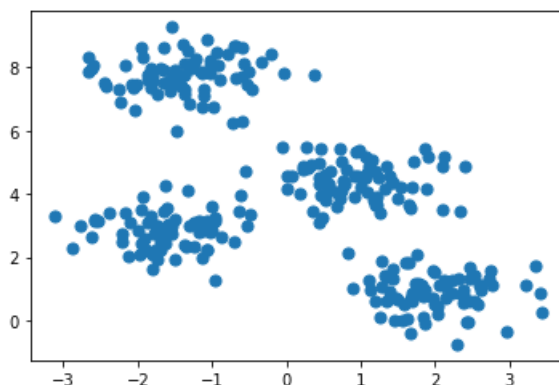
```
In [37]: from sklearn.datasets.samples_generator import make_blobs

X, y = make_blobs(n_samples=300, centers=4,
                  random_state=0, cluster_std=0.60)

plt.scatter(X[:, 0], X[:, 1], s=50);
```

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:144: FutureWarning: The sklearn.datasets.samples_generator module is deprecated in version 0.22 and will be removed in version 0.24. The corresponding classes / functions should instead be imported from sklearn.datasets. Anything that cannot be imported from sklearn.datasets is now part of the private API.

```
warnings.warn(message, FutureWarning)
```



- By eye, it is relatively easy to pick out the four clusters.
- If we were to perform an exhaustive search for the different segmentations of the data, however, the search space would be exponential in the number of points.
- Fortunately, there is a well-known Expectation Maximization (EM) procedure which scikit-learn implements, so that KMeans can be solved relatively quickly.

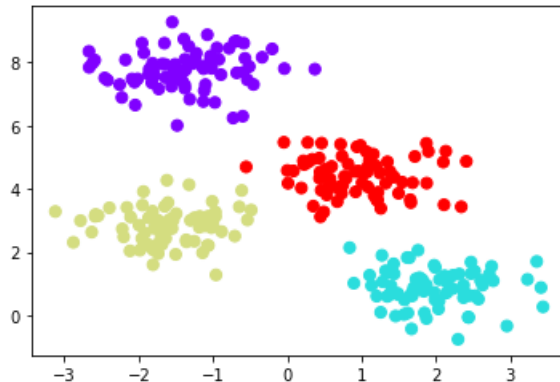
```
In [38]: from sklearn.cluster import KMeans

est = KMeans(4)

est.fit(X)

y_kmeans = est.predict(X)

plt.scatter(X[:, 0], X[:, 1], c=y_kmeans, s=50, cmap='rainbow');
```



```
In [39]: from sklearn.metrics import calinski_harabasz_score

calinski_harabasz_score(X, y_kmeans)
```

Out[39]: 1210.0899142587816

- The algorithm identifies the four clusters of points in a manner very similar to what we would do by eye!

Application of KMeans to Digits dataset

- Here we'll use K-Means to automatically cluster the data in 64 dimensions, and then look at the cluster centers to see what the algorithm has found.

```
In [40]: from sklearn.datasets import load_digits

digits = load_digits()
```

```
In [41]: est = KMeans(n_clusters=10)

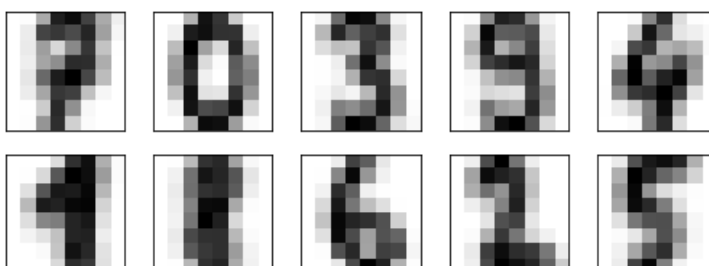
clusters = est.fit_predict(digits.data)

est.cluster_centers_.shape
```

Out[41]: (10, 64)

- We see ten clusters in 64 dimensions. Let's visualize each of these cluster centers to see what they represent:

```
In [42]: fig = plt.figure(figsize=(8, 3))
for i in range(10):
    ax = fig.add_subplot(2, 5, 1 + i, xticks=[], yticks=[])
    ax.imshow(est.cluster_centers_[i].reshape((8, 8)), cmap=plt.cm.binary)
```



```
In [43]: from sklearn.metrics import calinski_harabasz_score  
  
calinski_harabasz_score(digits.data, clusters)
```

Out[43]: 169.30606481277894

- We see that even without the labels, KMeans is able to find clusters whose means are recognizable digits (sorry to number 8)!

Principal Component Analysis

```
In [44]: import numpy as np  
from sklearn.decomposition import PCA  
X = np.array([[1, -1], [0, 1], [-1, 0]])  
print(X)  
print()  
pca = PCA()  
pca.fit(X)  
  
print(pca.components_)  
# Principal axes representing the directions of maximum variance in the data  
print()  
print(pca.explained_variance_ratio_)  
# Percentage of variance explained by each component  
print()  
X_new = pca.transform(X)  
print(X_new)  
  
[[ 1 -1]  
 [ 0  1]  
 [-1  0]]  
  
[[ 0.70710678 -0.70710678]  
 [ 0.70710678  0.70710678]]  
  
[0.75 0.25]  
  
[[ 1.41421356e+00 -3.33066907e-16]  
 [-7.07106781e-01  7.07106781e-01]  
 [-7.07106781e-01 -7.07106781e-01]]
```

```
In [45]: import numpy as np
from sklearn.decomposition import PCA
X = np.array([[ -1, -1, 5], [-2, -1, 8], [-3, -2, 6],
              [ 1, 1, 8], [ 2, 1, 5], [ 3, 2, 9]])

print(X)
print()
pca = PCA(n_components=2)
pca.fit(X)

print(pca.components_) # n_components x n_features
# Principal axes representing the directions of maximum variance in the data
print()
print(pca.explained_variance_ratio_)
# Percentage of variance explained by each component
print()
X_new = pca.transform(X)
print(X_new)
```

```
[[ -1 -1  5]
 [-2 -1  8]
 [-3 -2  6]
 [ 1  1  8]
 [ 2  1  5]
 [ 3  2  9]]

[[ 0.79554567  0.53157493  0.29074936]
 [ 0.31712316  0.04357793 -0.94738264]]

[0.76997386 0.22885582]

[[-1.86016108  1.37616707]
 [-1.78345867 -1.783104  ]
 [-3.69207799 -0.24903982]
 [ 1.66632818 -0.74457865]
 [ 1.58962577  2.41469243]
 [ 4.0797438  -1.01413703]]
```

- PCA on IRIS dataset

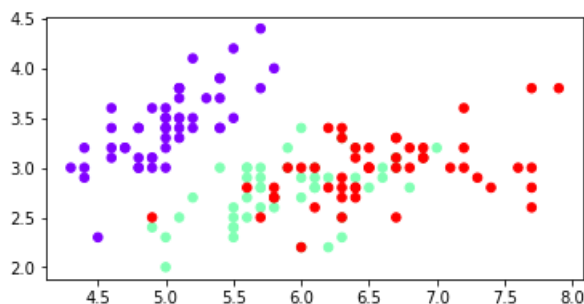
```
In [46]: import sklearn.datasets as ds
import matplotlib.pyplot as plt
%matplotlib inline

iris = ds.load_iris()
X = iris.data
y = iris.target
print(X.shape)

plt.figure(figsize=(6,3));
plt.scatter(X[:,0], X[:,1], c=y, s=30, cmap=plt.cm.rainbow)
```

(150, 4)

Out[46]: <matplotlib.collections.PathCollection at 0x7fbed2dde290>



```
In [47]: from sklearn.decomposition import PCA
pca = PCA()

X_pca = pca.fit_transform(X)

print(pca.components_)
print()
print(pca.explained_variance_ratio_)

[[ 0.36138659 -0.08452251  0.85667061  0.3582892 ]
 [ 0.65658877  0.73016143 -0.17337266 -0.07548102]
 [-0.58202985  0.59791083  0.07623608  0.54583143]
 [-0.31548719  0.3197231  0.47983899 -0.75365743]]

[0.92461872 0.05306648 0.01710261 0.00521218]
```

- Now display the same dataset, but in a new coordinate system (or equivalently, a linearly transformed version of the initial dataset).

```
In [48]: print(X_pca.shape)

(150, 4)
```

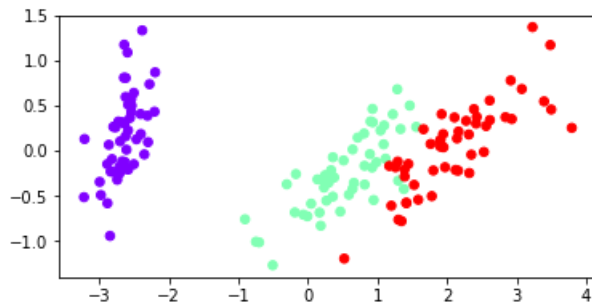
```
In [49]: X_pca[:,0]
# observe that dimensions 2 and 3 are almost zero values
# (change 0th column to 1st or 2nd or 3rd)
```

```
Out[49]: array([-2.68412563, -2.71414169, -2.88899057, -2.74534286, -2.72871654,
-2.28085963, -2.82053775, -2.62614497, -2.88638273, -2.6727558 ,
-2.50694709, -2.61275523, -2.78610927, -3.22380374, -2.64475039,
-2.38603903, -2.62352788, -2.64829671, -2.19982032, -2.5879864 ,
-2.31025622, -2.54370523, -3.21593942, -2.30273318, -2.35575405,
-2.50666891, -2.46882007, -2.56231991, -2.63953472, -2.63198939,
-2.58739848, -2.4099325 , -2.64886233, -2.59873675, -2.63692688,
-2.86624165, -2.62523805, -2.80068412, -2.98050204, -2.59000631,
-2.77010243, -2.84936871, -2.99740655, -2.40561449, -2.20948924,
-2.71445143, -2.53814826, -2.83946217, -2.54308575, -2.70335978,
 1.28482569,  0.93248853,  1.46430232,  0.18331772,  1.08810326,
 0.64166908,  1.09506066, -0.74912267,  1.04413183, -0.0087454 ,
-0.50784088,  0.51169856,  0.26497651,  0.98493451, -0.17392537,
 0.92786078,  0.66028376,  0.23610499,  0.94473373,  0.04522698,
 1.11628318,  0.35788842,  1.29818388,  0.92172892,  0.71485333,
 0.90017437,  1.33202444,  1.55780216,  0.81329065, -0.30558378,
-0.06812649, -0.18962247,  0.13642871,  1.38002644,  0.58800644,
 0.80685831,  1.22069088,  0.81509524,  0.24595768,  0.16641322,
 0.46480029,  0.8908152 ,  0.23054802, -0.70453176,  0.35698149,
 0.33193448,  0.37621565,  0.64257601, -0.90646986,  0.29900084,
 2.53119273,  1.41523588,  2.61667602,  1.97153105,  2.35000592,
 3.39703874,  0.52123224,  2.93258707,  2.32122882,  2.91675097,
 1.66177415,  1.80340195,  2.1655918 ,  1.34616358,  1.58592822,
 1.90445637,  1.94968906,  3.48705536,  3.79564542,  1.30079171,
 2.42781791,  1.19900111,  3.49992004,  1.38876613,  2.2754305 ,
 2.61409047,  1.25850816,  1.29113206,  2.12360872,  2.38800302,
 2.84167278,  3.23067366,  2.15943764,  1.44416124,  1.78129481,
 3.07649993,  2.14424331,  1.90509815,  1.16932634,  2.10761114,
 2.31415471,  1.9222678 ,  1.41523588,  2.56301338,  2.41874618,
 1.94410979,  1.52716661,  1.76434572,  1.90094161,  1.39018886])
```



```
In [50]: plt.figure(figsize=(6,3))
plt.scatter(X_pca[:,0], X_pca[:,1], c=y, s=30, cmap=plt.cm.rainbow)
```

```
Out[50]: <matplotlib.collections.PathCollection at 0x7fbecfd2e350>
```



- Points belonging to the same classes are now grouped together, even though the PCA estimator did not use the labels.
- The PCA was able to find a projection maximizing the variance, which corresponds here to a projection where the classes are well separated.

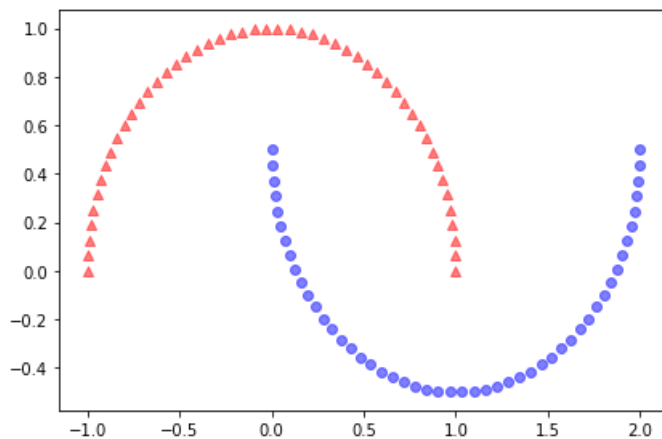
Kernel PCA

```
In [51]: from sklearn.decomposition import KernelPCA

X, y = ds.make_moons(n_samples=100)
# Synthetic Data: two interleaving half circles

plt.scatter(X[y == 0, 0], X[y == 0, 1], color='red', marker='^', alpha=0.5)
plt.scatter(X[y == 1, 0], X[y == 1, 1], color='blue', marker='o', alpha=0.5)

plt.tight_layout()
plt.show()
```



```
In [52]: # First, Let's try with PCA
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)

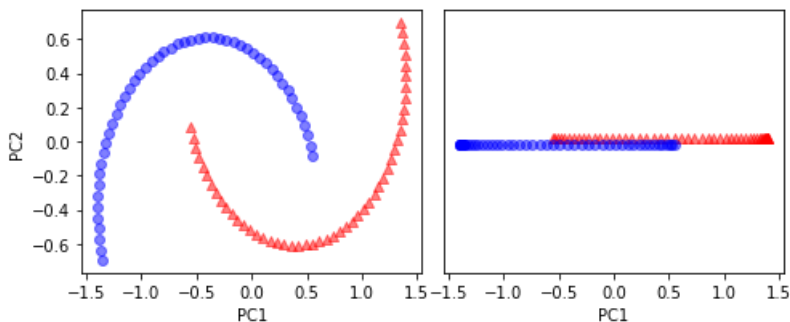
fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(7, 3))

ax[0].scatter(X_pca[y == 0, 0], X_pca[y == 0, 1],
              color='red', marker='^', alpha=0.5)
ax[0].scatter(X_pca[y == 1, 0], X_pca[y == 1, 1],
              color='blue', marker='o', alpha=0.5)

ax[1].scatter(X_pca[y == 0, 0], np.zeros((50, 1)) + 0.02,
              color='red', marker='^', alpha=0.5)
ax[1].scatter(X_pca[y == 1, 0], np.zeros((50, 1)) - 0.02,
              color='blue', marker='o', alpha=0.5)

ax[0].set_xlabel('PC1')
ax[0].set_ylabel('PC2')
ax[1].set_ylim([-1, 1])
ax[1].set_yticks([])
ax[1].set_xlabel('PC1')

plt.tight_layout()
plt.show()
```



```
In [53]: kpca = KernelPCA(n_components=2, kernel='rbf', gamma=15)
X_kpca = kpca.fit_transform(X)

plt.scatter(X_kpca[y == 0, 0], X_kpca[y == 0, 1],
            color='red', marker='^', alpha=0.5)
plt.scatter(X_kpca[y == 1, 0], X_kpca[y == 1, 1],
            color='blue', marker='o', alpha=0.5)

plt.xlabel('PC1')
plt.ylabel('PC2')
plt.tight_layout()
plt.show()
```

