

# Machine Learning with Python

## Machine Learning Workshop @ MPSTME

Instructor: Santosh Chapaneri

## Model Comparison with ROC Curves

```
In [1]: import numpy as np
import pandas as pd
import sklearn.model_selection as cv
```

```
In [2]: # Load the dataset with Pandas
train = pd.read_csv('Data/titanic_train.csv')

data = train[['Sex', 'Age', 'Pclass', 'Survived']].copy()
data['Sex'] = data['Sex'] == 'female'
data = data.dropna()
data.head()
```

Out[2]:

	Sex	Age	Pclass	Survived
0	False	22.0	3	0
1	True	38.0	1	1
2	True	26.0	3	1
3	True	35.0	1	1
4	False	35.0	3	0

```
In [3]: # Create X and Y
data_np = data.astype(np.int32).values
X = data_np[:, :-1] # Features
y = data_np[:, -1] # Target (survived or not)
```

```
In [4]: # Split the dataset
X_train, X_test, y_train, y_test = cv.train_test_split(X, y, test_size=0.25)
```

```
In [5]: from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import BernoulliNB
from sklearn import svm
from sklearn.linear_model import Perceptron
from sklearn.neural_network import MLPClassifier

from sklearn import metrics
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
%matplotlib inline
```

## Logistic Regression Classifier

```
In [6]: ### Logistic Regression
logreg = LogisticRegression(solver='lbfgs')
logreg.fit(X_train, y_train)

y_pred_log = logreg.predict(X_test)

confmat = confusion_matrix(y_test, y_pred_log)
confmat
```

```
Out[6]: array([[78, 22],
               [25, 54]], dtype=int64)
```

```
In [7]: accuracy_LOG = accuracy_score(y_test, y_pred_log)
print('Accuracy of LogReg:', accuracy_LOG)
print()

from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_log))
```

Accuracy of LogReg: 0.7374301675977654

	precision	recall	f1-score	support
0	0.76	0.78	0.77	100
1	0.71	0.68	0.70	79
avg / total	0.74	0.74	0.74	179


```
In [8]: ### ROC Curve for Logistic Regression
LOG_roc_auc = roc_auc_score(y_test, logreg.predict(X_test))
print('AUC of LogReg:', LOG_roc_auc)

fpr_LOG, tpr_LOG, thresholds_LOG = roc_curve(y_test,
                                              logreg.predict_proba(X_test)[:,:1])

def plot_roc_curve(fpr, tpr, label=None):
    plt.plot(fpr, tpr, linewidth=2, label=label)
    plt.plot([0,1],[0,1], 'k--')
    plt.axis([0,1,0,1])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')

plot_roc_curve(fpr_LOG, tpr_LOG)
plt.show()
```

AUC of LogReg: 0.7317721518987342



## KNN Classifier

```
In [9]: knn_clf = KNeighborsClassifier(n_neighbors=5)
knn_clf.fit(X_train, y_train)
y_pred_KNN = knn_clf.predict(X_test)
print(classification_report(y_test, y_pred_KNN))


accuracy_KNN = accuracy_score(y_test, y_pred_KNN)
print('Accuracy of KNN:', accuracy_KNN)

KNN_roc_auc = roc_auc_score(y_test, knn_clf.predict(X_test))
print('AUC of KNN:', KNN_roc_auc)

fpr_KNN, tpr_KNN, thresholds_KNN = roc_curve(y_test,
                                              knn_clf.predict_proba(X_test)[:,-1])
plot_roc_curve(fpr_KNN, tpr_KNN)
plt.show()
```

	precision	recall	f1-score	support
0	0.72	0.76	0.74	100
1	0.68	0.63	0.65	79
avg / total	0.70	0.70	0.70	179

Accuracy of KNN: 0.7039106145251397  
AUC of KNN: 0.6964556962025317



## Naive Bayes Classifier

```
In [10]: ### Gaussian Naive Bayes Algorithm
bnb = BernoulliNB()
bnb.fit(X_train, y_train)
y_pred_bnb = bnb.predict(X_test)
print(classification_report(y_test, y_pred_bnb))


accuracy_BNB = accuracy_score(y_test, y_pred_bnb)
print('Accuracy of Naive Bayes:', accuracy_BNB)

BNB_roc_auc = roc_auc_score(y_test, bnb.predict(X_test))
print('AUC of Naive Bayes:', BNB_roc_auc)

fpr_BNB, tpr_BNB, thresholds_BNB = roc_curve(y_test,
                                             bnb.predict_proba(X_test)[: ,1])
plot_roc_curve(fpr_BNB, tpr_BNB)
plt.show()
```

	precision	recall	f1-score	support
0	0.76	0.76	0.76	100
1	0.70	0.70	0.70	79
avg / total	0.73	0.73	0.73	179

Accuracy of Naive Bayes: 0.7318435754189944  
AUC of Naive Bayes: 0.7281012658227849



## SVM

```
In [11]: SVM_clf = svm.SVC(kernel='linear',probability=True)
SVM_clf.fit(X_train, y_train)
y_pred_SVM = SVM_clf.predict(X_test)
print(classification_report(y_test, y_pred_SVM))


accuracy_SVM = accuracy_score(y_test, y_pred_SVM)
print('Accuracy of SVM:', accuracy_SVM)

SVM_roc_auc = roc_auc_score(y_test, SVM_clf.predict(X_test))
print('AUC of SVM:', SVM_roc_auc)

fpr_SVM, tpr_SVM, thresholds_SVM = roc_curve(y_test,
                                             SVM_clf.predict_proba(X_test)[:,:1])
plot_roc_curve(fpr_SVM, tpr_SVM)
plt.show()
```

	precision	recall	f1-score	support
0	0.75	0.76	0.76	100
1	0.69	0.68	0.69	79
avg / total	0.73	0.73	0.73	179

Accuracy of SVM: 0.7262569832402235  
AUC of SVM: 0.7217721518987342



```
In [12]: # SVM with hyper-paramete tuning
# Set the parameters by cross-validation
parameters = [{'kernel': ['rbf'],
                        'gamma': [1e-4, 1e-3, 0.01, 0.1, 0.2, 0.5],
                        'C': [1, 10, 100, 1000]},
               {'kernel': ['linear'], 'C': [1, 10, 100, 1000]}]

# Perform Grid Search
print("# Tuning hyper-parameters")
from sklearn.model_selection import GridSearchCV
clf = GridSearchCV(svm.SVC(probability=True), parameters, cv=5, iid=True)

# Fit
clf.fit(X_train, y_train)

# Predict on test data
y_pred_SVM_CV = clf.predict(X_test)

# Show the classification report
print(classification_report(y_test, y_pred_SVM_CV))

accuracy_SVM_CV = accuracy_score(y_test, y_pred_SVM_CV)
print('Accuracy of SVM:', accuracy_SVM_CV)

SVM_CV_roc_auc = roc_auc_score(y_test, clf.predict(X_test))
print('AUC of SVM:', SVM_CV_roc_auc)


fpr_SVM_CV, tpr_SVM_CV, thresholds_SVM_CV = roc_curve(y_test,
                                                         clf.predict_proba(X_test)[:,-1])
plot_roc_curve(fpr_SVM_CV, tpr_SVM_CV)
plt.show()
```

```
# Tuning hyper-parameters
              precision    recall  f1-score   support

         0         0.74      0.81      0.78       100
         1         0.73      0.65      0.68        79

avg / total         0.74      0.74      0.74       179
```

```
Accuracy of SVM: 0.7374301675977654
AUC of SVM: 0.7277848101265822
```



## Perceptron

```
In [17]: perc_clf = Perceptron(alpha=0.0001, class_weight=None, eta0=1.0,
                             fit_intercept=True, max_iter=1000, n_jobs=1,
                             penalty=None, random_state=0, shuffle=True,
                             verbose=0, warm_start=False)

perc_clf.fit(X_train, y_train)
y_pred_PERC = perc_clf.predict(X_test)
print(classification_report(y_test, y_pred_PERC))

accuracy_PERC = accuracy_score(y_test, y_pred_PERC)
print('Accuracy of Perceptron:', accuracy_PERC)
```

	precision	recall	f1-score	support
0	0.71	0.05	0.09	100
1	0.45	0.97	0.61	79
avg / total	0.60	0.46	0.32	179

Accuracy of Perceptron: 0.4581005586592179

## Multi-layer Perceptron

```
In [18]: mlp = MLPClassifier(hidden_layer_sizes=(13,13,13),max_iter=500)
mlp.fit(X_train,y_train)
y_pred_MLP = mlp.predict(X_test)
print(classification_report(y_test, y_pred_MLP))

accuracy_MLP = accuracy_score(y_test, y_pred_MLP)
print('Accuracy of MLP:', accuracy_MLP)
```

	precision	recall	f1-score	support
0	0.57	0.99	0.73	100
1	0.83	0.06	0.12	79
avg / total	0.69	0.58	0.46	179

Accuracy of MLP: 0.5810055865921788

```
In [15]: # Snapshot
print ("Logistic Regression Accuracy : " ,accuracy_LOG)
print ("KNN Accuracy : " ,accuracy_KNN)
print ("Naive Bayes Accuracy : " ,accuracy_BNB)
print ("SVM Accuracy : " ,accuracy_SVM)
print ("SVM_CV Accuracy : " ,accuracy_SVM_CV)
print ("Perceptron Accuracy : " ,accuracy_PERC)
print ("MLP Accuracy : " ,accuracy_MLP)
```

```
Logistic Regression Accuracy : 0.7374301675977654
KNN Accuracy : 0.7039106145251397
Naive Bayes Accuracy : 0.7318435754189944
SVM Accuracy : 0.7262569832402235
SVM_CV Accuracy : 0.7374301675977654
Perceptron Accuracy : 0.4581005586592179
MLP Accuracy : 0.5083798882681564
```

In [19]: *# Plot ALL ROC curves in one plot*

```
plt.figure()
plt.plot(fpr_LOG, tpr_LOG,
         label='LogReg Model (area = %0.2f)' % LOG_roc_auc)
plt.plot(fpr_KNN, tpr_KNN,
         label='KNN Model (area = %0.2f)' % KNN_roc_auc)
plt.plot(fpr_BNB, tpr_BNB,
         label='NB Model (area = %0.2f)' % BNB_roc_auc)
plt.plot(fpr_SVM, tpr_SVM,
         label='SVM Model (area = %0.2f)' % SVM_roc_auc)
plt.plot(fpr_SVM_CV, tpr_SVM_CV,
         label='SVM_CV Model (area = %0.2f)' % SVM_CV_roc_auc)

plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.show()
```

