

ML Hands-on Workshop @ Elec, SFIT

Instructor: Santosh Chapaneri

Jan 2022

```
In [25]: import numpy as np
import pandas as pd
import sklearn.svm as svm
import matplotlib.pyplot as plt
# %matplotlib inline
```

Support Vector Machines (SVM)

- We generate 2D points and assign a binary label according to a linear operation on the coordinates.

```
In [26]: X = np.random.randn(200, 2)
y = (X[:, 0] + X[:, 1]) > 1 # Imagine an AND function

X[:10], y[:10]
```

```
Out[26]: (array([[ 1.10182890e+00,  1.63997803e-01],
 [ 5.54692088e-01,  5.12114836e-01],
 [-1.90611261e-02, -9.13936739e-01],
 [ 5.40503603e-01, -1.44962686e-01],
 [-9.66361265e-01, -4.05920547e-01],
 [ 1.02819928e-01,  5.06715832e-03],
 [-8.09612152e-01,  2.68775885e+00],
 [-7.71194000e-01, -9.14522368e-01],
 [-8.14323032e-01,  2.11823106e-03],
 [-2.31353109e-01, -4.90825246e-01]]),
 array([ True,  True, False, False, False, False,  True, False, False,
        False]))
```

- Fit a linear Support Vector Classifier (SVC)

```
In [27]: est = svm.LinearSVC()
est.fit(X, y)
```

```
Out[27]: LinearSVC()
```

```

In [28]: # Generate a grid in the square  $[-3,3]^2$ 
xx, yy = np.meshgrid(np.linspace(-3, 3, 500),
                     np.linspace(-3, 3, 500))

# This function takes a SVM estimator as input
def plot_decision_function(est, X, y):
    # We evaluate the decision function on the grid.
    Z = est.decision_function(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    cmap = plt.cm.Blues

    # Display the decision function on the grid
    plt.figure(figsize=(5,5));
    plt.imshow(Z,
               extent=(xx.min(), xx.max(), yy.min(), yy.max()),
               aspect='auto', origin='lower', cmap=cmap)

    # Display the boundaries
    plt.contour(xx, yy, Z, levels=[0], linewidths=2, colors='k')

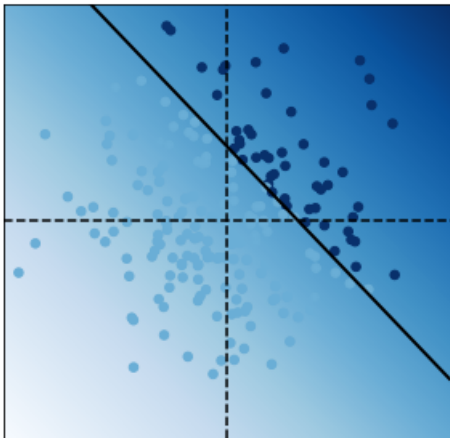
    # Display the points with their true labels
    plt.scatter(X[:, 0], X[:, 1], s=30, c=.5+.5*y, lw=1,
               cmap=cmap, vmin=0, vmax=1);
    plt.axhline(0, color='k', ls='--')
    plt.axvline(0, color='k', ls='--')
    plt.xticks(())
    plt.yticks(())
    plt.axis([-3, 3, -3, 3])

```

```

In [29]: plot_decision_function(est, X, y);
#plt.title("Linearly separable, linear SVC")

```




The linear SVC tried to separate the points with a line and it did a good job.

- We now modify the labels with a XOR function.
- A point's label is 1 if the coordinates have different signs. This classification is not linearly separable. Therefore, a linear SVC fails completely.

```
In [30]: y = np.logical_xor(X[:, 0] > 0, X[:, 1] > 0)

est2 = svm.LinearSVC()
est2.fit(X, y)


plot_decision_function(est2, X, y)
#plt.title("XOR, Linear SVC")
```



- It is possible to use non-linear SVCs by using non-linear kernels.
- Kernels specify a non-linear transformation of the points into a higher-dimensional space. Transformed points in this space are assumed to be more linearly separable, although they are not necessarily in the original space.
- By default, the SVC classifier in scikit-learn uses the Radial Basis Function (RBF) kernel.

```
In [31]: est3 = svm.SVC()
est3.fit(X, y)


plot_decision_function(est3, X, y)
```



```
In [32]: from sklearn.model_selection import GridSearchCV

est4 = GridSearchCV(svm.SVC(),
                    {'C': np.logspace(-3., 3., 10),
                     'gamma': np.logspace(-3., 3., 10)}, cv=5);
est4.fit(X, y)
#print("Score: {0:.3f}".format(cv.cross_val_score(est4, X, y).mean()))

plot_decision_function(est4.best_estimator_, X, y)
```



Applying SVM on IRIS dataset

```
In [33]: from sklearn import datasets

iris = datasets.load_iris()
X = iris.data
y = iris.target
```

- Let's do **hyper-parameter tuning**
- Use **5-fold cross validation** to perform grid search to calculate optimal hyper-parameters

```
In [34]: from sklearn.model_selection import GridSearchCV
import sklearn.model_selection as cv
from sklearn.metrics import classification_report

# Split the dataset
X_train, X_test, y_train, y_test = cv.train_test_split(X, y, test_size=0.25)
```

```
In [35]: # Set the parameters by cross-validation
parameters = [{'kernel': ['rbf'],
                      'gamma': [1e-4, 1e-3, 0.01, 0.1, 0.2, 0.5],
                      'C': [1, 10, 100, 1000]},
               {'kernel': ['linear'], 'C': [1, 10, 100, 1000]}]

print("# Tuning hyper-parameters")
print()

clf = GridSearchCV(svm.SVC(decision_function_shape='ovr'), parameters, cv=5)
clf.fit(X_train, y_train)

# Tuning hyper-parameters
```

```
Out[35]: GridSearchCV(cv=5, estimator=SVC(),
                    param_grid=[{'C': [1, 10, 100, 1000],
                                'gamma': [0.0001, 0.001, 0.01, 0.1, 0.2, 0.5],
                                'kernel': ['rbf']},
                               {'C': [1, 10, 100, 1000], 'kernel': ['linear']}])
```

```
In [36]: print("Best parameters set found on training set:")  
print()  
print(clf.best_params_)
```

Best parameters set found on training set:

```
{'C': 1, 'gamma': 0.5, 'kernel': 'rbf'}
```

```
In [37]: y_true, y_pred = y_test, clf.predict(X_test)  
print(classification_report(y_true, y_pred))  
# The support is the number of occurrences of each class in y_true
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	16
1	0.77	1.00	0.87	10
2	1.00	0.75	0.86	12
accuracy			0.92	38
macro avg	0.92	0.92	0.91	38
weighted avg	0.94	0.92	0.92	38

```
In [38]: from sklearn.metrics import accuracy_score  
accuracy_score(y_true, y_pred)
```

Out[38]: 0.9210526315789473