

Natural Language Processing (NLP)

RNN/LSTM for Text Processing

Instructor: Santosh Chapaneri

Sep 2021

- The IMDB dataset contains 25,000 highly-polar movie reviews (good or bad) for training and the same amount again for testing.
- The problem is to determine whether a given movie review has a positive or negative sentiment.
- The words have been replaced by integers that indicate the ordered frequency of each word in the dataset. The sentences in each review are therefore comprised of a sequence of integers.

Word Embedding

- We will map each word onto a 32 length real valued vector. We will also limit the total number of words that we are interested in modeling to the 5000 most frequent words, and zero out the rest. Finally, the sequence length (number of words) in each review varies, so we will constrain each review to be 500 words, truncating long reviews and pad the shorter reviews with zero values.

```
In [ ]: import numpy as np
        from keras.datasets import imdb
        from keras.models import Sequential
        from keras.layers import Dense, LSTM, SimpleRNN, GRU
        from keras.layers import LSTM
        from keras.layers.embeddings import Embedding
        from keras.preprocessing import sequence
        # fix random seed for reproducibility
        np.random.seed(2021)
```

- We are constraining the dataset to the top 5,000 words. We also split the dataset into train (50%) and test (50%) sets.

```
In [ ]: # Load the dataset but only keep the top n words, zero the rest
        top_words = 5000

        (X_train, y_train), (X_test, y_test) = imdb.load_data(num_words = top_words)
```

```
<string>:6: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray
/usr/local/lib/python3.7/dist-packages/keras/datasets/imdb.py:155: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray
    x_train, y_train = np.array(xs[:idx]), np.array(labels[:idx])
/usr/local/lib/python3.7/dist-packages/keras/datasets/imdb.py:156: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray
    x_test, y_test = np.array(xs[idx:]), np.array(labels[idx:])
```

```
In [ ]: X_train.shape
```

```
Out[ ]: (25000,)
```

```
In [ ]: len(X_train[2]) # vary number to see length of different samples
```

```
Out[ ]: 141
```

```
In [ ]: X_train[2][:10]
```

```
Out[ ]: [1, 14, 47, 8, 30, 31, 7, 4, 249, 108]
```

- We need to truncate and pad the input sequences so that they are all the same length for modeling.
- The model will learn the zero values carry no information so indeed the sequences are not the same length in terms of content, but same length vectors is required to perform the computation in Keras.

```
In [ ]: import keras
```

```
# Retrieve the word index file mapping words to indices
word_index = keras.datasets.imdb.get_word_index()

# Reverse the word index to obtain a dict mapping indices to words
inverted_word_index = dict((i, word) for (word, i) in word_index.items())

# Decode the first sequence in the dataset
idx = 25
decoded_sequence = " ".join(inverted_word_index[i] for i in X_train[idx])
decoded_sequence
```

```
Out[ ]: "the as it is time usual basis must has is small whole for there is works oh and most all low in a
nd they be martial and developed in long an friendly br appeal br of great this is playing and br
and and to recently in also of clearly br is save br specially past mixed or actually french mysel
f and there is copy editing like book else damon show and to it look so and finds and br and or is
dislike more he something br budget what's better of and this were and film and dave to and early
around get of every that it girl each in perfect man second some br of and film as you not like dr
ew that it see is you in own have is again older they is hell certainly way this and"
```

```
In [ ]: # truncate and pad input sequences
```

```
max_review_length = 500
```

```
X_train_pad = sequence.pad_sequences(X_train, maxlen = max_review_length, padding = 'post')
```

```
X_test_pad = sequence.pad_sequences(X_test, maxlen = max_review_length, padding = 'post')
```

```
In [ ]: X_train_pad.shape
```

```
Out[ ]: (25000, 500)
```



```
In [ ]: from keras.layers import SimpleRNN

# create the model
embedding_veclen = 32

model = Sequential()
model.add(Embedding(top_words, embedding_veclen, input_length = max_review_length))
model.add(SimpleRNN(100))
model.add(Dense(1, activation='sigmoid'))

model.compile(loss = 'binary_crossentropy',
              optimizer = 'adam',
              metrics = ['accuracy'])

model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
embedding_2 (Embedding)	(None, 500, 32)	160000

simple_rnn_2 (SimpleRNN)	(None, 100)	13300

dense_2 (Dense)	(None, 1)	101
=====		
Total params: 173,401		
Trainable params: 173,401		
Non-trainable params: 0		

```
In [ ]: model.fit(X_tr, y_tr,
                validation_data=(X_val, y_val),
                epochs=3, batch_size=64)
```

```
Epoch 1/3
16/16 [=====] - 8s 452ms/step - loss: 0.6984 - accuracy: 0.5152 - val_loss: 0.6922 - val_accuracy: 0.5240
Epoch 2/3
16/16 [=====] - 6s 389ms/step - loss: 0.6944 - accuracy: 0.5108 - val_loss: 0.6930 - val_accuracy: 0.5160
Epoch 3/3
16/16 [=====] - 6s 404ms/step - loss: 0.6814 - accuracy: 0.5676 - val_loss: 0.6938 - val_accuracy: 0.5000
```

```
Out[ ]: <keras.callbacks.History at 0x7ff916dd5fd0>
```

```
In [ ]: # Final evaluation of the model

scores = model.evaluate(X_test, y_test)
scores
```

```
16/16 [=====] - 1s 33ms/step - loss: 0.6937 - accuracy: 0.5180
```

```
Out[ ]: [0.6936662793159485, 0.5180000066757202]
```

• Stacked RNN

```
In [ ]: model_st_rnn = Sequential()
model_st_rnn.add(Embedding(top_words, embedding_veclen, input_length = max_review_length))
model_st_rnn.add(SimpleRNN(100, return_sequences = True))
model_st_rnn.add(SimpleRNN(100, return_sequences = False))
model_st_rnn.add(Dense(1, activation='sigmoid'))
model_st_rnn.compile(loss = 'binary_crossentropy',
                    optimizer = 'adam',
                    metrics = ['accuracy'])
model_st_rnn.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
embedding_5 (Embedding)	(None, 500, 32)	160000
simple_rnn_7 (SimpleRNN)	(None, 500, 100)	13300
simple_rnn_8 (SimpleRNN)	(None, 100)	20100
dense_5 (Dense)	(None, 1)	101
Total params: 193,501		
Trainable params: 193,501		
Non-trainable params: 0		

```
In [ ]: model_st_rnn.fit(X_tr, y_tr,
                        validation_data=(X_val, y_val),
                        epochs=3, batch_size=64)
```

```
Epoch 1/3
16/16 [=====] - 24s 834ms/step - loss: 0.7176 - accuracy: 0.5054 - val_loss: 0.7161 - val_accuracy: 0.5240
Epoch 2/3
16/16 [=====] - 12s 777ms/step - loss: 0.6998 - accuracy: 0.5234 - val_loss: 0.6947 - val_accuracy: 0.4880
Epoch 3/3
16/16 [=====] - 12s 781ms/step - loss: 0.6884 - accuracy: 0.5279 - val_loss: 0.6938 - val_accuracy: 0.4900
```

```
Out[ ]: <keras.callbacks.History at 0x7ff91402e690>
```

```
In [ ]: # Final evaluation of the model

scores = model_st_rnn.evaluate(X_test, y_test)
scores
```

```
16/16 [=====] - 1s 64ms/step - loss: 0.6921 - accuracy: 0.5260
```

```
Out[ ]: [0.6921359300613403, 0.5260000228881836]
```

LSTM

```
In [ ]: model_lstm = Sequential()
model_lstm.add(Embedding(top_words, embedding_veclen, input_length = max_review_length))
model_lstm.add(LSTM(100, return_sequences = False))
model_lstm.add(Dense(1, activation='sigmoid'))
model_lstm.compile(loss = 'binary_crossentropy',
                    optimizer = 'adam',
                    metrics = ['accuracy'])
model_lstm.summary()
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
embedding_6 (Embedding)	(None, 500, 32)	160000
lstm (LSTM)	(None, 100)	53200
dense_6 (Dense)	(None, 1)	101
Total params: 213,301		
Trainable params: 213,301		
Non-trainable params: 0		

```
In [ ]: model_lstm.fit(X_tr, y_tr,
                      validation_data=(X_val, y_val),
                      epochs=3, batch_size=64)
```

Epoch 1/3
 16/16 [=====] - 7s 70ms/step - loss: 0.6931 - accuracy: 0.5136 - val_loss: 0.6937 - val_accuracy: 0.4940
 Epoch 2/3
 16/16 [=====] - 1s 36ms/step - loss: 0.6928 - accuracy: 0.4918 - val_loss: 0.6930 - val_accuracy: 0.4920
 Epoch 3/3
 16/16 [=====] - 1s 34ms/step - loss: 0.6920 - accuracy: 0.5478 - val_loss: 0.6938 - val_accuracy: 0.4920

Out[]: <keras.callbacks.History at 0x7ff90e656490>

```
In [ ]: # Final evaluation of the model

scores = model_lstm.evaluate(X_test, y_test)
scores
```

16/16 [=====] - 0s 16ms/step - loss: 0.6918 - accuracy: 0.5360

Out[]: [0.6918402314186096, 0.5360000133514404]

Deep LSTM

```
In [ ]: model_st_lstm = Sequential()
model_st_lstm.add(Embedding(top_words, embedding_vec_len, input_length = max_review_length))
model_st_lstm.add(LSTM(100, return_sequences = True))
model_st_lstm.add(LSTM(100, return_sequences = False))
model_st_lstm.add(Dense(1, activation='sigmoid'))
model_st_lstm.compile(loss = 'binary_crossentropy',
                      optimizer = 'adam',
                      metrics = ['accuracy'])
model_st_lstm.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
embedding_7 (Embedding)	(None, 500, 32)	160000
lstm_1 (LSTM)	(None, 500, 100)	53200
lstm_2 (LSTM)	(None, 100)	80400
dense_7 (Dense)	(None, 1)	101
Total params: 293,701		
Trainable params: 293,701		
Non-trainable params: 0		

```
In [ ]: model_st_lstm.fit(X_tr, y_tr,
                        validation_data=(X_val, y_val),
                        epochs=3, batch_size=64)
```

Epoch 1/3
 16/16 [=====] - 5s 114ms/step - loss: 0.6939 - accuracy: 0.5086 - val_loss: 0.6951 - val_accuracy: 0.4760
 Epoch 2/3
 16/16 [=====] - 1s 58ms/step - loss: 0.6936 - accuracy: 0.5029 - val_loss: 0.6939 - val_accuracy: 0.4760
 Epoch 3/3
 16/16 [=====] - 1s 61ms/step - loss: 0.6924 - accuracy: 0.5253 - val_loss: 0.6936 - val_accuracy: 0.4880

Out[]: <keras.callbacks.History at 0x7ff9198ffd50>

```
In [ ]: # Final evaluation of the model

scores = model_st_lstm.evaluate(X_test, y_test)
scores
```

16/16 [=====] - 0s 26ms/step - loss: 0.6925 - accuracy: 0.5380

Out[]: [0.6924556493759155, 0.5379999876022339]

GRU

```
In [ ]: model_gru = Sequential()
model_gru.add(Embedding(top_words, embedding_veclen, input_length = max_review_length))
model_gru.add(GRU(100, return_sequences = False))
model_gru.add(Dense(1, activation='sigmoid'))
model_gru.compile(loss = 'binary_crossentropy',
                  optimizer = 'adam',
                  metrics = ['accuracy'])
model_gru.summary()
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
embedding_8 (Embedding)	(None, 500, 32)	160000
gru (GRU)	(None, 100)	40200
dense_8 (Dense)	(None, 1)	101
Total params: 200,301		
Trainable params: 200,301		
Non-trainable params: 0		

```
In [ ]: model_gru.fit(X_tr, y_tr,
                    validation_data=(X_val, y_val),
                    epochs=3, batch_size=64)
```

Epoch 1/3

16/16 [=====] - 3s 69ms/step - loss: 0.6936 - accuracy: 0.4780 - val_loss: 0.6941 - val_accuracy: 0.4780

Epoch 2/3

16/16 [=====] - 0s 31ms/step - loss: 0.6925 - accuracy: 0.5223 - val_loss: 0.6942 - val_accuracy: 0.4820

Epoch 3/3

16/16 [=====] - 0s 28ms/step - loss: 0.6909 - accuracy: 0.5617 - val_loss: 0.6931 - val_accuracy: 0.4840

Out[]: <keras.callbacks.History at 0x7ff91510c550>

```
In [ ]: # Final evaluation of the model

scores = model_gru.evaluate(X_test, y_test)
scores
```

16/16 [=====] - 0s 17ms/step - loss: 0.6926 - accuracy: 0.5420

Out[]: [0.6925719380378723, 0.5419999957084656]